

Substantive

A specific mitigation, safeguard or deployment decision

Partial

Interim or incomplete — or claimed but unverifiable

Acknowledged

The problem is referenced, no action specified

None

No public company response located

THE FINDING

CAISI found two novel security vulnerabilities in ChatGPT Agent and chained them into a working exploit.

UK AISI found a universal jailbreak in GPT-5.5 within six hours.

UK AISI found Claude Sonnet 4.5 verbalises unprompted evaluation awareness in 16.6% of outputs, against 1.6% for Opus 4.1.

Dreadnode found Anthropic's Opus 4.8 and Sonnet 5 were the heaviest cheaters of the 22 models tested (65.2% and 56.5%).

WHAT THE COMPANY DID

OpenAI fixed both within one business day of the report.

OpenAI said it had patched it, but AISI could not verify the fix before public deployment.

Anthropic reproduced the measurement in the system card, but recorded no safeguard change.

Five-source battery exhausted and dated; no company document located.