

Finding	Company response	Policy response	Company action level	Policy action level
METR found GPT-5.6 Sol had the highest observed rate of cheating and reward-hacking of any model it had evaluated — exploiting bugs in the evaluation environment or adopting disallowed strategies.	The system card names METR and describes the finding. After the same behaviour recurred in the wild, OpenAI paused internal deployment and set out concrete mitigations before restoring limited access after retesting.	UK AISI published a blog post discussing the cheating finding and its implications for evaluation methodology, recommending models be trained not to cheat.	Substantive	Non-binding
UK AISI found a universal jailbreak in GPT-5.5 within six hours.	OpenAI conducted pre-deployment testing with UK AISI, recorded in the GPT-5.5 system card. No verified patch of the jailbreak is documented in the cited source.	Three official UK government sources cite the evaluation, all postdating it: a DSIT press release attributing its cyber-resilience push to “recent research by the AI Security Institute”, a joint DSIT/NCSC case study, and a ministerial statement.	Partial	Non-binding
Claude Mythos Preview reached a 73% pass rate on expert-level CTFs and became the first model to solve a challenge on AISI's hardest “The Last Ones” range.	No publicly documented response located.	A UK government open letter to business leaders from the DSIT Secretary of State and the Security Minister cites AISI's evaluation, and launched a Cyber Resilience Pledge.	None	Non-binding