
Do third-party frontier AI evaluations matter?

Anonymous Author(s)

Affiliation

Address

email

Abstract

Government AI institutes and independent evaluators study frontier AI systems to identify, evaluate, and better understand potential risks. Yet identifying and evaluating a risk does not itself mitigate the risk. We examine whether published third-party AI evaluation findings concerning a named frontier system are acted upon. The analysis covers publicly documented findings from government AI institutes and independent evaluation organisations through August 2026. We trace downstream responses through three channels: company responses and model updates; parliamentary and regulatory action; and documented media and academic coverage. The corpus contains 1,146 findings, of which 233 concerned empirical results about named systems for which a company response could reasonably be expected. Of the 190 that met our predefined significant-risk threshold, 152 (80%) lacked a specific, publicly documented company response meeting our proportionality criteria. Together, these findings suggest that, based on publicly available information, there is only limited reason to trust that third-party evaluations of frontier AI systems consistently lead to proportionate, attributable, and verifiable responses. We compare the accountability structures currently in place for frontier AI systems with those established by oversight bodies in other industries, including food, drugs, energy, transportation, and finance. We conclude that formal mechanisms for access, response, remediation, verification, and follow-up inspired by governance in these other sectors could increase the consistency with which AI evaluations translate to meaningful action.

1 Introduction

Frontier AI system evaluations have become a prominent component of AI governance. Government AI institutes and independent evaluators test advanced systems for dangerous capabilities, safeguard failures, and other risks, often with access and technical expertise unavailable to the general public. These evaluations can provide evidence about an AI system’s risk beyond the boundaries of the developers’ own disclosures [1]. However, identifying a risk is not the same as reducing it. An evaluation can document a serious vulnerability without prompting a developer response, such as a deployment delay, additional review, or verification of remediation. Therefore, the practical value of the evaluation depends partly on what happens after the finding is communicated or published.

Across the 190 significant-risk findings for which a company response could be assessed, 152 (80%) had no publicly documented response meeting our proportionality criteria, while only 38 (20%) received a proportionate response (Table 1). The shortfall appeared among findings reported by both government institutes and independent evaluators and was substantially greater for post-deployment than pre-deployment evaluations.

An absence of a publicly documented response does not establish that no action occurred or that the finding was unimportant. Preventive, private, and unattributed company actions may leave no public trace, so our results describe the public record rather than companies’ internal conduct.

39 Finally, we compare the frontier-AI evaluation ecosystem with oversight institutions in pharmaceu-
 40 tical regulation, nuclear safety, transportation, finance, and related domains. These institutions are
 41 not direct templates for AI governance, but they illustrate formal mechanisms for access, response,
 42 remediation, verification, escalation, and continuing follow-up.

Table 1: Outcomes among the 190 C1 (significant risk) Tier A (trackable evaluation) findings.

Outcome	Count	% of C1	What it means
Accountability gap (no action)	114	60%	Alarming finding concerning a named company, with no publicly documented company response.
Under-response (partial / acknowledged)	38	20%	The company acknowledged the finding or described a response, but no complete and verifiable mitigation was documented.
Proportionate	38	20%	A specific, documented response meeting the proportionality criteria; attribution is assessed separately.
Gap total (rows 1 + 2)	152	80%	Either no response or an under-response; no documented company response met the proportionality standard.

44 2 Background

45 2.1 Government AI institutes

46 Government AI institutes emerged as technical public bodies focused on evaluating the safety of
 47 advanced AI systems. The first institutes were established in the United Kingdom, the United States,
 48 and Japan during 2023 and 2024. Their initial mandates broadly combined safety evaluation, re-
 49 search, standards development, and international cooperation [2].

50 These institutes have since developed in different directions. In February 2025, the UK AI Safety
 51 Institute was renamed the AI Security Institute and its remit was refocused on risks with security
 52 implications, including cyberattacks, chemical and biological threats, fraud, and other forms of
 53 criminal misuse [3]. It remains within the Department for Science, Innovation and Technology
 54 and focuses on research and evaluation rather than regulatory enforcement. In June 2025, the US
 55 AI Safety Institute was re-established as the Center for AI Standards and Innovation (US CAISI)
 56 within the National Institute of Standards and Technology. Its work combines model evaluation
 57 and national-security assessment with measurement science, guidelines, and voluntary standards
 58 development [4]. Japan’s AI Safety Institute, based within the Information-technology Promotion
 59 Agency, primarily develops evaluation methods, guidance, standards, and testing infrastructure [5].
 60 France’s INESIA, launched in 2025, is not a new standalone agency but a coordinated network of
 61 existing public bodies working on systemic-risk analysis, model evaluation, and support for the
 62 implementation of AI regulation [6].

63 Government AI institutes therefore no longer represent a single institutional model. They vary in
 64 focus, structure, access arrangements, and connection to standards or regulatory processes. Most
 65 findings examined in this paper are technical outputs that may inform developers and policymakers
 66 but do not themselves compel a company response or remediation. We use “government AI insti-
 67 tutes” as an umbrella term for the public bodies in our sampling frame that conduct or coordinate
 68 external evaluations of advanced AI systems. We distinguish them from independent evaluators
 69 while applying the same public-record response framework to findings reported by both groups.

70 2.2 Independent AI system evaluators and third-party auditing

71 Independent evaluators are defined as non-government organisations (including non-profit organi-
 72 sations, for-profit evaluation firms, and academic research groups) that are institutionally separate
 73 from the developer whose system they evaluate and that conduct and publicly report frontier-AI
 74 evaluations. They include generalist evaluators such as METR and Apollo Research and domain-
 75 specific organisations such as SecureBio. Independent evaluators also test frontier systems for dan-
 76 gerous capabilities, safeguard failures, and other risks, but differ from government institutes in insti-
 77 tutional status, funding, access arrangements, and public responsibilities. Their access may depend

on voluntary developer cooperation, contractual arrangements, application programming interfaces, or publicly available models, allowing some to conduct pre-deployment evaluations while limiting others to post-deployment testing. At the same time, independent evaluators can provide specialised expertise and scrutiny outside both government and frontier-model developers. Recent work therefore treats evaluator independence, system access, methodological rigour, reporting arrangements, and audit standards as important dimensions of third-party auditing [7–10]. This paper treats government institutes and independent evaluators as distinct institutional categories while applying the same public-record response framework to eligible findings from both.

3 Methodology

3.1 Data collection

We constructed the sampling frame by identifying 46 government bodies and independent evaluators that conducted frontier-AI evaluations and publicly reported their methods and results. We screened 6,684 publications and identified 457 reports containing sufficiently documented evaluations of advanced AI systems. The earliest included report was published on March 17, 2023, and the public-record cutoff was August 29, 2026. From these reports, we coded 1,146 distinct findings. The corpus includes adverse findings about named systems, findings about anonymised systems, reassuring or null results, capability trends, and methodology or governance findings. Section 3.2 places them in three tiers. Only the 233 Tier A (trackable evaluation) findings—adverse findings about a named model or developer—enter the response and proportionality analyses. The sampling frame and the 46 reporting-institution labels, including separate labels for joint evaluations, are listed in Appendices A and D.

The unit of analysis is an individual finding rather than a report. A finding is one discrete, evidence-backed, independently codeable claim concerning one accountable company. It must be asserted by the evaluator and supported by a measured result, observed behaviour, or documented process fact. Opinions, recommendations, announcements, and plans are excluded. Results are combined when they concern the same company, describe the same issue, and would reasonably call for a common response. They are coded separately when they describe substantively different issues or would reasonably call for different responses. Results concerning different companies are always coded as separate findings, even when the same report identifies the same problem in multiple companies’ models. Results concerning multiple models from one company may be combined when they document the same issue and would reasonably call for the same response. Each finding receives a unique Finding ID, while findings drawn from the same report retain a common Report ID.

The corpus is restricted to evidence available in the public record. Findings communicated privately but not subsequently disclosed cannot be observed. We also excluded non-English-language outputs where a sufficiently reliable English version was unavailable. The resulting corpus should therefore be interpreted as a structured collection of publicly accessible, predominantly English-language evidence, not as a complete census of every evaluation performed by the included institutions.

3.1.1 Tooling and human oversight

Corpus construction was AI-assisted. Automated screening and extraction identified candidate publications and findings, while AI agents searched separately for Channel A (company responses), Channel B (policy uptake), and Channel C (public coverage). Two authors independently reviewed each retained finding and its coding, including eligibility, tier, severity, response evidence, and Action Level; the agreed human judgment determined the final coding. Appendix B gives the search instructions, source classes, documented effort, and limitations.

3.2 Defining the accountability set

We assigned each finding to one of three mutually exclusive tiers based on whether it reported an adverse or safety-relevant result and permitted developer-specific response analysis.

Tier A (Trackable evaluation findings): Response-trackable adverse findings. A finding entered Tier A (Trackable evaluation findings) when it: (1) reported a specific adverse or safety-relevant result; (2) identified a named model, system, or developer rather than only an anonymised placeholder; and (3) concerned an issue within the developer’s capacity to address, such that a company response

could reasonably be assessed. These 233 findings constitute the accountability set and are the only findings included in the Action Level and proportionality analyses.

Tier B (Non-trackable evaluation findings): Empirical findings that did not meet at least one Tier A condition. These include anonymised-system results, reassuring or null results, bare scores or rankings without a concerning threshold, non-frontier systems, capability trends, inconclusive results, and findings for which a company response could not reasonably be assessed.

Tier C (Non-empirical findings): Methodology, framework, governance or process, tooling, milestone, and other findings that did not report empirical model behaviour or capability.

All three tiers remain in the descriptive corpus. Only Tier A (Trackable evaluation findings) enters the company-response and proportionality analyses.

3.3 Severity of risk classification

A finding was classified as C1 (significant risk) when the reported evidence met at least one significant-risk threshold across eight risk domains: chemical, biological, radiological, and nuclear (CBRN) uplift; offensive cyber capability; autonomy, self-replication, or AI research-and-development (AI-R&D) automation; persuasion or societal harm at scale; deliberate deception or misalignment; deployed-safeguard failure; compromised evaluation integrity; or acute individual harm. Findings could be assigned to more than one domain. A finding was classified as C2 (low risk) when the reported evidence did not meet any of these significant-risk thresholds. This does not necessarily mean that the finding was unimportant or that no response was warranted. The eight domains draw on two sources: published developer frameworks and additional risk criteria developed for this study. Domains D1–D3—CBRN uplift, offensive cyber capability, and autonomy including AI-R&D automation and self-replication—correspond to dangerous-capability thresholds that frontier developers publish for their own systems in the Anthropic Responsible Scaling Policy [11], the OpenAI Preparedness Framework [12], and the Google DeepMind Frontier Safety Framework [13]. Domains D4 and D5—persuasion or societal harm at scale and demonstrated deceptive misalignment—correspond to categories that those frameworks treat less uniformly. Domains D6–D8—deployed-safeguard failure, compromised evaluation integrity, and acute individual harm—are additional criteria developed for this study. They describe failures observed in the evaluation literature that the developers’ capability thresholds do not cover.

The majority vote of the three models (Claude Sonnet 5, GPT-5.5, and Gemini 3.1 Pro) provided the initial severity classification. Two authors then independently reviewed every finding and its supporting evidence, including the individual model votes and rationales. Where warranted, the authors overrode the model majority; the agreed human-reviewed label was the final severity classification used in the analysis. Individual model votes and model-generated reasons were retained so that model disagreements remained observable. Appendix C describes the annotation procedure; the complete prompt, boundary rules, and output format are available with the supplemental study materials.

Across the full corpus, 1,032 of 1,146 findings (90.1%) received a unanimous 3–0 classification, while 114 (9.9%) were determined by a 2–1 vote. Among the 233 Tier A (trackable evaluation) findings, 194 (83.3%) received a unanimous classification and 39 (16.7%) were determined by a 2–1 vote. The final corpus contained 340 C1 (significant risk) and 806 C2 (low risk) findings. Within Tier A (trackable evaluation), 190 findings were classified C1 (significant risk) and 43 were classified C2 (low risk).

3.4 Response evidence and outcome coding

3.4.1 Evidence channels

Table 2: Evidence channels searched separately through August 29, 2026. Channels B and C do not count as a company response.

Channel	Evidence examined	Role in the analysis
Channel A (company response)	Company statements, system cards, and safety or deployment pages.	Attribution and Action Level.
Channel B (policy uptake)	Official parliamentary, legislative, regulatory, and government records.	Policy Level.
Channel C (public coverage)	Independent media coverage, academic citations, and notable public discussion.	External visibility.

175 The resulting public-record evidence was coded into four study-defined measures: Attribution, Ac-
 176 tion Level, Policy Level, and Proportionality.

177 3.4.2 Attribution

178 Attribution records whether an identified response was publicly connected to the relevant finding.
 179 (a) Explicit attribution: The responding company directly referenced the finding, report, evaluation
 180 result, or evaluating institution, or stated that its response was informed by that evidence. (b) No
 181 explicit attribution: A qualifying response was identified, but the company did not publicly state
 182 a connection to the finding. (c) No response located: No qualifying company response was iden-
 183 tified, so attribution does not arise. Attribution is assessed only where a company response exists.
 184 Cases with no identified response are recorded under the third category rather than left blank, so
 185 that a searched-and-empty result is distinguishable from an uncoded one. Attribution measures the
 186 company's public explanation and does not independently establish causation.

187 3.4.3 Action Level

188 Action Level records the strength of the publicly documented company response identified through
 189 Channel A (Company response). (a) None: No public response could be identified through the
 190 search by the cutoff date. (b) Acknowledged: The company recognised the finding or underlying
 191 problem but specified no action. (c) Partial: The company documented an action that addressed
 192 only part of the identified problem or was expressly described as interim, limited, or incomplete.
 193 (d) Substantive: The company documented a specific mitigation, model change, safeguard, access
 194 restriction, or deployment decision that directly addressed the identified problem. Action Level
 195 measures the content of the public response. It does not establish whether the action was successfully
 196 implemented or effective in reducing the risk. Only public primary documentation issued by the
 197 responding company was admitted as Channel A evidence; evaluator statements, news coverage,
 198 and other third-party sources did not themselves qualify as company responses.

199 3.4.4 Policy Level

200 Policy Level records whether the finding received documented uptake through Channel B (policy
 201 uptake) and whether it created an enforceable requirement. The categories are No policy uptake
 202 identified, Non-binding policy-related uptake, and Binding policy action.

203

SEVERITY	COMPANY RESPONSE LEVEL			
	Substantive	Partial	Acknowledged	None
C1 <i>significant risk</i>	PROPORTIONATE	UNDER-RESPONSE	UNDER-RESPONSE	ACCOUNTABILITY GAP
C2 <i>low risk</i>	PROPORTIONATE	PROPORTIONATE	UNDER-RESPONSE	ACCOUNTABILITY GAP

Figure 1: Proportionality coding rule. Rows show severity; columns show company-response level; cells show the resulting outcome. The matrix defines the rule, not observed results.

204 3.4.5 Proportionality

205 Proportionality combines two independently assigned classifications: finding severity and company
 206 Action Level. In Figure 1, the rows show Severity (C1 significant risk or C2 low risk), and the
 207 columns show the documented company-response level. The matrix shows the severity-dependent
 208 rule used to classify every Tier A (trackable evaluation) finding. Equivalently, a C1 (significant risk)
 209 finding requires a Substantive response to meet the proportionality standard, whereas a C2 (low
 210 risk) finding requires at least a Partial response. Acknowledgment without action never meets the

standard, and the absence of a company response always produces a no-action gap. This is a study-defined measure of documented response proportionality. It evaluates whether the publicly described response met the minimum Action Level associated with the finding’s severity classification. It does not establish that the response was effective, that remediation was successfully implemented, or that the underlying risk was eliminated.

4 Results

We first report company-response outcomes among the 190 C1 (significant risk) findings in Tier A (trackable evaluation findings), which constitute the paper’s primary accountability analysis. We then examine differences by evaluation-access type, public coverage of findings without company responses, and the extent to which reported company responses could be independently verified.

4.1 Company responses to significant-risk findings

Restricted to C1 (significant risk) findings—the comparison relevant to the paper’s central accountability question—152 of 190 findings (80%; Wilson 95% CI 74–85%) received either no publicly documented company response or an inadequate one. This combined figure, rather than the 62.2% no-response rate across all 233 Tier A (trackable evaluation) findings, is the paper’s central result. Because it is calculated only over C1 findings, it is not mechanically affected by the number of C2 (low risk) findings included in the corpus. Of the 76 C1 findings with a located company response, 64 had explicit attribution and 12 did not. In Channel B (policy uptake), 41 C1 findings received non-binding policy uptake, one received binding action, and 148 had no identified policy uptake. Robustness checks and subgroup tests are reported in Appendix F.

4.2 Outcomes by evaluation access type

Among the C1 (significant risk) Tier A (trackable evaluation) findings classified as either pre- or post-deployment, the substantive-response rate was 55% for the 44 pre-deployment findings and 7% for the 131 post-deployment findings. Pre-deployment findings were therefore roughly eight times as likely to receive a substantive response. No-response rates also differed substantially: 6.8% for pre-deployment findings and 79.4% for post-deployment findings. A two-proportion test gave $z = -8.54$, $p < 0.001$, and the difference remained significant after adjustment for clustering within reports ($p < 0.001$; Appendix F.2 and Figure F.2). These comparisons describe an association and should not be interpreted causally.

This descriptive association is consistent with pre-deployment access facilitating company action, although access type is correlated with evaluator, disclosure arrangements, domain, and other institutional differences. One plausible mechanism is visible in the case-level record: pre-deployment findings are communicated while the deployment decision window remains open and are frequently published under coordinated-disclosure arrangements. Forty-nine of the 88 recorded response lags were zero days, meaning that the finding and company response were published together. This remains a plausible interpretation rather than a mechanism established by the comparison itself.

4.3 Public coverage of findings without company responses

Of the 114 C1 (significant risk) Tier A (trackable evaluation) findings classified as Accountability gap (no action), 91 (80%) have at least one mainstream press item or academic citation recorded. Including notable public discussion, 92 (81%) have evidence in at least one Channel C (public coverage) field. Thus, the absence of a publicly documented company response was not generally confined to findings with no observable external coverage. This evidence does not establish that the relevant company saw the coverage, but it shows that many no-response findings were publicly visible beyond the original evaluation report. The channel-by-channel breakdown and the inconclusive traction comparison are reported in Appendix F.2 and Table F.2.

4.4 Independent verification of company responses

A company response does not necessarily permit independent verification. During its pre-deployment evaluation of GPT-5.5, UK AISI identified a universal jailbreak within six hours of testing [14]. OpenAI subsequently reported several changes to its safeguard stack. However, a

260 configuration issue in the version supplied to UK AISI meant the institute could not verify the effective-
261 ness of the final configuration before deployment; the cited source does not document a verified
262 patch of the jailbreak.

263 This case illustrates the distinction captured by the Partial category: a company may document a
264 response while the external evaluator remains unable to verify whether the final mitigation addressed
265 the identified vulnerability.

266 The limitation this case exposes is structural rather than incidental. Channel A (company response)
267 admits only the responding company’s own primary documentation (Section 3.4), so every one of
268 the 76 C1 (significant risk) findings that drew a response is evidenced by that company’s account of
269 itself. That is a property of the coding rule, not a finding. The question the rule cannot answer is
270 whether a second party ever confirmed that the described action was taken, or that it worked.

271 Reading the response record for each of those 76 findings, we identify one case in which an eval-
272 uator is documented as having re-tested a revised safeguard and endorsed it: the US Center for AI
273 Standards and Innovation’s (US CAISI’s) examination of Anthropic’s rebuilt cyber classifier, which
274 occurred during the export-control process and before redeployment (Section 5.1). We identify one
275 case in which the evaluator states that it could not verify the claimed fix—the GPT-5.5 jailbreak
276 described above. For the remaining 74, the public record contains the company’s description and
277 nothing independent against which to check it.

278 The count of independently verified remediations should be read as a lower bound. It reflects what
279 the cited sources describe; a verification exercise that took place but was not reported would not
280 appear. Four further candidates were examined and excluded: collaboration on pre-deployment
281 capability testing is not confirmation of a remedy, a capability designation is not a verified mitigation,
282 a developer confirming an attack is not a developer demonstrating a fix, and a company’s assessment
283 of its own model using its own tooling is not independent. The Partial category captures the subset
284 where the shortfall is visible at coding time—21 of the 190 C1 findings (11%)—but the broader
285 position is that proportionality throughout this paper is assessed against what companies state they
286 did, not against remediation anyone else has checked. Appendix G, especially Tables G.1 and G.2,
287 gives worked examples of this boundary.

288 **Limitations.** These results describe the public record, not unreported company action, and do not
289 establish that an evaluation caused a response. Findings also cluster within reports, so subgroup
290 comparisons use report-level adjustments.

291 5 Understanding the accountability gap

292 Section 4 documents the distribution of publicly observable responses but does not, by itself, ex-
293 plain why those outcomes occurred. This section examines two forms of follow-through: company
294 remediation in two illustrative cases and policy-record responses to findings in the corpus. These
295 analyses help identify circumstances associated with observable action, but they are neither causal
296 estimates nor exhaustive explanations of the accountability gap. Section 6 separately compares the
297 follow-up mechanisms available to government AI institutes with those used by established over-
298 sight institutions.

299 To illustrate how evaluation findings can be followed by documented remediation despite the broader
300 accountability gap, we examine two cases with an unusually clear public sequence from technical
301 evidence to developer response. Both are included in the corpus and coded as receiving a Sub-
302 stantive response. In the Fable 5 and Mythos 5 case, the triggering Amazon report falls outside
303 the evaluator-scope sampling frame, but the later US CAISI safeguard-testing and redeployment
304 record is included as finding USCAISI-2026-06-CYB2. The cases contrast legally consequential
305 action through an existing government authority with voluntary remediation through iterative pre-
306 deployment collaboration. They are illustrative rather than representative.

307 5.1 Enforcement through an existing legal instrument: Fable 5 and Mythos 5

308 Anthropic launched Claude Fable 5 on June 9, 2026. According to Anthropic, the US government
309 applied export controls to Fable 5 and Mythos 5 on June 12, requiring the company to prevent access

310 by foreign nationals inside and outside the United States. But Anthropic suspended access to both
311 models worldwide.

312 Anthropic reported that the directive followed government awareness of an Amazon research report
313 describing a technique for bypassing Fable 5’s safeguards. The technique caused the model to iden-
314 tify software vulnerabilities and, in one instance, produce code demonstrating how one vulnerability
315 could be exploited. Anthropic subsequently trained an updated safety classifier and reported that it
316 blocked the identified technique in more than 99% of cases. Anthropic also reported that US CAISI
317 tested both the previous and updated safeguards and considered them “extraordinarily strong”; this
318 does not establish that US CAISI independently produced the 99% estimate. The controls were
319 lifted on June 30. Fable 5 returned to global availability on July 1, while Mythos 5 was initially
320 restored only to selected US organisations. [15].

321 The triggering Amazon report is not a corpus finding because it does not satisfy the evaluator-scope
322 criteria. The later US CAISI safeguard-testing and redeployment record is included as finding
323 USCAISI-2026-06-CYB2, and is the one corpus response for which the public record states that
324 an evaluator tested the revised safeguards. The public record therefore documents a comparatively
325 complete accountability sequence: a technical problem was reported, a binding legal instrument im-
326 posed restrictions, the company implemented a mitigation, and a government evaluator examined
327 the revised safeguards before the restrictions were lifted.

328 The case does not show that the US federal government possessed or exercised deployment authority.
329 The legally consequential action came through the Department of Commerce’s export-control pow-
330 ers, while US CAISI contributed technical testing within the resulting process. It therefore illustrates
331 how evaluation expertise can support remediation when another institution supplies an existing legal
332 pathway.

333 5.2 Collaborative pre-deployment remediation: GPT-5.6

334 OpenAI’s system card describes an iterative pre-deployment process rather than a single evaluation
335 event. UK AISI received early access to GPT-5.6 Sol and extensive grey-box access to its safe-
336 guard system. Across several testing rounds, UK AISI reported universal cyber-domain jailbreaks,
337 including attacks that enabled long-form agentic work involving vulnerability discovery and exploit
338 development. OpenAI states that it improved the system in response to UK AISI’s findings and, by
339 launch, had reproduced and mitigated the specific jailbreaks UK AISI reported. The system card
340 does not claim that the broader vulnerability class was eliminated: UK AISI expected further red-
341 teaming to identify similar jailbreaks, and OpenAI committed to continued testing with the institute.
342 [16].

343 We code the response as Substantive because OpenAI publicly attributes implemented pre-
344 deployment mitigation to the jailbreaks reported by UK AISI. This classification refers specifically
345 to mitigation of the reported attacks; it does not imply that GPT-5.6 had become resistant to all simi-
346 lar jailbreaks. Unlike the Fable 5 case, no binding government order was involved. The documented
347 response occurred through repeated testing, private disclosure, mitigation, and retesting within an
348 existing evaluator–developer collaboration.

349 Together, the two cases illustrate distinct pathways from technical evidence to remediation. In the
350 Fable 5 case, an existing legal instrument imposed consequences, while government evaluation sup-
351 ported verification of the revised safeguards. In the GPT-5.6 case, pre-deployment access and an
352 established evaluator–developer collaboration were followed by documented voluntary mitigation.
353 Neither case establishes that publication of an evaluation finding, by itself, reliably triggers follow-
354 up.

355 6 Institutional mechanisms for reducing the accountability gap

356 The accountability problem examined in this paper is not unique to AI governance. Established
357 oversight systems distinguish between producing technically credible findings, obtaining a response,
358 verifying corrective action, and possessing authority to compel or escalate action. We compare the
359 UK AI Security Institute (UK AISI) and US Center for AI Standards and Innovation (US CAISI)
360 with seven established oversight bodies in the United States and the United Kingdom, focusing on
361 formal mechanisms for information access, designated recipients, response, remediation or escala-

tion, verification, and continuing follow-up. These bodies differ substantially in mandate, subject matter, and enforcement authority and are not presented as direct institutional templates for frontier AI governance. The comparison instead identifies mechanisms through which technical findings can remain actionable and publicly visible after an investigation or evaluation has concluded.

Detailed histories of the National Transportation Safety Board (NTSB) and Public Company Accounting Oversight Board (PCAOB), together with the wider institutional comparison, are provided in Appendix H and Table H.1 [17–21].

Neither precedent provides a ready-made model for frontier AI governance. The NTSB cannot force recipients to adopt its recommendations, while the PCAOB regulates auditors rather than directly compelling audited companies to remedy every adverse finding. Nor do these precedents establish that AI evaluation institutions have failed in the same ways. They instead separate two institutional questions that are often conflated: whether a system provides a formal pathway from technical findings to documented follow-up, and whether the evaluators themselves are subject to independent supervision. Applied to frontier AI evaluation, these precedents motivate asking who should respond to an adverse finding, who tracks and evaluates that response, who verifies claimed remediation, and who—if anyone—oversees the evaluators.

The comparative lesson is therefore institutional rather than sector-specific. Government frontier-AI evaluators need not themselves become comprehensive regulators for their findings to enter a more systematic accountability process. Technical evaluation could remain institutionally separate while findings are connected to designated recipients, response expectations or obligations, public status tracking, procedures for reviewing claimed remediation, and defined pathways for escalation or referral. These mechanisms would not guarantee that every finding is remediated or that every response reduces real-world risk, but they would make follow-through more consistent, attributable, verifiable, and publicly inspectable.

7 Conclusion

Government AI institutes and independent evaluators increasingly provide evidence about frontier systems that would otherwise remain available only through developer self-reporting. The findings in this paper do not question the value of that technical work. They examine whether the public record shows that alarming findings concerning named systems were followed by responses proportionate to the identified risk.

Closing this public accountability gap does not necessarily require converting technical institutes into general-purpose AI regulators. Companies participating in government evaluations could publish finding-level responses stating whether they accept, dispute, plan to address, or have remediated a finding. Institutes could maintain a public response-status record, subject to legitimate security and confidentiality constraints, and distinguish claimed remediation from independently verified remediation. Evaluation agreements could also establish response timelines and specify what status information may be disclosed.

The public record currently provides limited grounds for confidence that government and independent frontier-AI evaluations consistently lead to substantive improvements in deployed-system safety. Creating a visible and repeatable pathway from finding to response, follow-up, and verification would make that confidence more warranted, and make the contribution of the evaluation ecosystem more publicly accessible.

References

- [1] T. Shevlane, S. Farquhar, B. Garfinkel, et al. (2023). Model evaluation for extreme risks. *arXiv preprint arXiv:2305.15324*. <https://arxiv.org/abs/2305.15324>.
- [2] R. Araujo, K. Fort, and O. Guest (2024). Understanding the first wave of AI Safety Institutes: Characteristics, functions, and challenges. *arXiv preprint arXiv:2410.09219*. <https://arxiv.org/abs/2410.09219>.
- [3] UK Government (2025). *Tackling AI security risks to unleash growth and deliver Plan for Change*. Department for Science, Innovation and Technology, February 14, 2025. <https://www.gov.uk/government/news/tackling-ai-security-risks-to-unleash-growth-and-deliver-plan-for-change>.
- [4] Center for AI Standards and Innovation (2026). *Center for AI Standards and Innovation (CAISI)*. National Institute of Standards and Technology. <https://www.nist.gov/caisi>.
- [5] Japan AI Safety Institute (2026). *Japan AI Safety Institute (J-AISI)*. Information-technology Promotion Agency. https://aisi.go.jp/assets/pdf/20260708_AISI_en.pdf.
- [6] French Government (2025). Le Gouvernement annonce la création de l’Institut national pour l’évaluation et la sécurité de l’intelligence artificielle (INESIA). *Ministère de l’Économie, des Finances et de la Souveraineté industrielle et numérique*, January 31, 2025. French Government announcement.
- [7] G. Falco, B. Shneiderman, J. Badger, et al. (2021). Governing AI safety through independent audits. *Nature Machine Intelligence*, 3, 566–571. <https://doi.org/10.1038/s42256-021-00370-7>.
- [8] S. Costanza-Chock, E. Harvey, I. D. Raji, M. Czernuszenko, and J. Buolamwini (2023). Who audits the auditors? Recommendations from a field scan of the algorithmic auditing ecosystem. *arXiv preprint arXiv:2310.02521*. <https://arxiv.org/abs/2310.02521>.
- [9] M. Brundage, N. Dreksler, A. Homewood, et al. (2026). Frontier AI auditing: Toward rigorous third-party assessment of safety and security practices at leading AI companies. *arXiv preprint arXiv:2601.11699*. <https://arxiv.org/abs/2601.11699>.
- [10] A. Birhane, R. Steed, V. Ojewale, B. Vecchione, and I. D. Raji (2024). AI auditing: The broken bus on the road to AI accountability. *arXiv preprint arXiv:2401.14462*. <https://arxiv.org/abs/2401.14462>.
- [11] Anthropic (2026). *Responsible Scaling Policy, Version 3.4*. Effective July 8, 2026. <https://www.anthropic.com/responsible-scaling-policy>.
- [12] OpenAI (2025). *Preparedness Framework, Version 2*. Last updated April 15, 2025. <https://cdn.openai.com/pdf/18a02b5d-6b67-4cec-ab64-68cdfbddebcd/preparedness-framework-v2.pdf>.
- [13] Google DeepMind (2026). *Frontier Safety Framework, Version 3.1*. Published April 17, 2026. <https://deepmind.google/frontier-safety/>.
- [14] OpenAI (2026). *GPT-5.5 System Card: External Evaluations for Cyber Capabilities—UK AISI*. <https://deploymentsafety.openai.com/gpt-5-5/external-evaluations-for-cyber-capabilities---uk-aisi>.
- [15] Anthropic (2026). *Redeploying Fable 5*. <https://www.anthropic.com/news/redeploying-fable-5>.
- [16] OpenAI (2026). *GPT-5.6 System Card*. <https://deploymentsafety.openai.com/gpt-5-6>.
- [17] U.S. Congress (2026). *49 U.S.C. Section 1135: Secretary of Transportation’s Responses to Safety Recommendations*. <https://uscode.house.gov/view.xhtml?req=%28title%3A49+section%3A1135+edition%3Aprelim%29>.
- [18] National Transportation Safety Board (2024). *Responding to Our Safety Recommendations*. <https://www.ntsb.gov/investigations/Pages/RecipientTipRecommendations.aspx>.
- [19] U.S. General Accounting Office (2002). *Protecting the Public Interest: Selected Governance, Regulatory Oversight, Auditing, Accounting, and Financial Reporting Issues*. Report GAO-02-483T. <https://www.gao.gov/products/gao-02-483t>.
- [20] U.S. Securities and Exchange Commission (2003). *Order Regarding Section 101(d) of the Sarbanes–Oxley Act of 2002*. Release Nos. 33-8223 and 34-47746. <https://www.sec.gov/rules-regulations/2003/04/order-regarding-section-101d-sarbanes-oxley-act-2002>.
- [21] Public Company Accounting Oversight Board (2013). *Staff Guidance Concerning the Remediation Process*. https://pcaobus.org/oversight/inspections/remediation/remediation_process.

- [22] U.S. Food and Drug Administration (n.d.). *About Warning and Close-Out Letters*. Accessed September 7, 2026. <https://www.fda.gov/inspections-compliance-enforcement-and-criminal-investigations/warning-letters/about-warning-and-close-out-letters>.
- [23] U.S. Nuclear Regulatory Commission (n.d.). *Inspection*. Accessed September 7, 2026. <https://www.nrc.gov/about-nrc/regulatory/safety-oversight>.
- [24] Office for Nuclear Regulation (n.d.). *How we regulate*. Accessed September 7, 2026. <https://www.onr.org.uk/our-work/how-we-regulate>.
- [25] Financial Conduct Authority (2026). *Our approach to supervision*. Updated March 26, 2026. <https://www.fca.org.uk/publications/corporate-documents/our-approach-to-supervision>.
- [26] UK Civil Aviation Authority (2023). *145.B.300 Oversight principles*. <https://regulatorylibrary.caa.co.uk/1321-2014/Content/Document%20Structure/02%20145/2%20Regs/145.B.300%20Oversight%20principles.htm>.
- [27] UK Civil Aviation Authority (n.d.). *Enforcement and prosecutions*. Accessed September 7, 2026. <https://www.caa.co.uk/about-us/the-civil-aviation-authority/enforcement/enforcement-and-prosecutions/>.
- [28] I. D. Raji, A. Smart, R. N. White, et al. (2020). Closing the AI accountability gap: Defining an end-to-end framework for internal algorithmic auditing. In *Proceedings of the 2020 Conference on Fairness, Accountability, and Transparency*, pages 33–44. <https://doi.org/10.1145/3351095.3372873>.
- [29] J. Mökander, J. Schuett, H. R. Kirk, and L. Floridi (2024). Auditing large language models: A three-layered approach. *AI and Ethics*, 4, 1085–1115. First published online in 2023. <https://doi.org/10.1007/s43681-023-00289-2>.
- [30] M. Anderljung, E. T. Smith, J. O’Brien, et al. (2023). Towards publicly accountable frontier LLMs: Building an external scrutiny ecosystem under the ASPIRE framework. *arXiv preprint arXiv:2311.14711*. <https://arxiv.org/abs/2311.14711>.
- [31] I. D. Raji, P. Xu, C. Honigsberg, and D. E. Ho (2022). Outsider oversight: Designing a third party audit ecosystem for AI governance. *arXiv preprint arXiv:2206.04737*. <https://arxiv.org/abs/2206.04737>.
- [32] L. Staufer, M. Yang, A. Reuel, and S. Casper (2025). Audit cards: Contextualizing AI evaluations. *arXiv preprint arXiv:2504.13839*. <https://arxiv.org/abs/2504.13839>.
- [33] Y. Bengio et al. (2026). *International AI Safety Report 2026*. <https://arxiv.org/abs/2602.21012>.
- [34] I. D. Raji and J. Buolamwini (2019). Actionable auditing: Investigating the impact of publicly naming biased performance results of commercial AI products. In *Proceedings of the 2019 AAAI/ACM Conference on AI, Ethics, and Society*, pages 429–435. <https://doi.org/10.1145/3306618.3314244>.
- [35] J. Wang, K. Huang, K. Klyman, and R. Bommasani (2025). Do AI companies make good on voluntary commitments to the White House? *arXiv preprint arXiv:2508.08345*. <https://arxiv.org/abs/2508.08345>.
- [36] T. Rost (2026). From disclosure to self-referential opacity: Six dimensions of strain in current AI governance. *arXiv preprint arXiv:2604.14070*. <https://arxiv.org/abs/2604.14070>.

Appendix

A Sampling frame and source eligibility

The appendices follow the research process: sampling and search (Appendices A–B), coding (Appendix C), corpus and response evidence (Appendices D–E), robustness and worked examples (Appendices F–G), institutional context and prior work (Appendices H–I), and reproducibility (Appendix J).

A.1 Detailed eligibility and source scope

The 1,146 evaluation findings included outputs from the UK AI Security Institute, the US Center for AI Standards and Innovation, national and joint evaluation initiatives involving Australia, France, Japan, Singapore, and South Korea, and exercises conducted through the International Network for Advanced AI Measurement, Evaluation and Science. The sampling frame also included third-party evaluators such as METR, SecureBio, Apollo Research, Palisade Research, and other non-profit, academic, and commercial evaluation organisations.

The initial screening roster was derived from three sources: national institutes and initiatives identified through the International Network for Advanced AI Measurement, Evaluation and Science; members of the AI Evaluators Forum plus evaluator organisations already represented in the developing corpus; and frontier-developer release sites, searched only for material explicitly attributed to an external evaluator. Organisations later found to be co-author affiliations rather than publishing evaluators, together with national institutes that produced no qualifying report, remain in Table A.1 and the screening ledger as zero-yield records.

Third-party outputs were eligible whether or not the work was conducted with a government institute, provided that the public material identified the evaluator and tested system and documented the methods and results sufficiently to extract a finding. Relevant sources included formal evaluation reports, research publications, institute progress reports containing evaluation results, joint-testing publications, and company model or system cards that explicitly reported or linked to an external evaluation. Company materials describing only internal evaluations were not treated as independent evaluation reports.

Table A.1 lists the organisations whose publications were screened to build the corpus, with the number of publications screened and reports included.

Table A.1: The 46 organisations in the sampling frame.

Organisation screened	Publications screened	Reports included
RAND Corporation	759	6
Kenya (Ministry of ICT and Digital Economy - network member)	628	0
Allen Institute for AI (Ai2)	547	0
Holistic AI	431	9
Google DeepMind	377	0
Scale AI	324	81
Anthropic	302	28
Cisco (Robust Intelligence / Foundation AI)	287	13
Thorn	270	0
UC Berkeley	269	0
OpenAI	166	20
UK AI Security Institute (UK AISI)	160	67
Center for AI Safety (CAIS)	156	13
EPFL	155	0
Redwood Research	137	15
Tsinghua University CoAI Group	136	3
FAR.AI	117	19
Korea AI Safety Institute (K-AISI)	105	4
Japan AI Safety Institute (J-AISI)	100	3
METR	93	49
UL Research Institutes - Digital Safety Research Institute (DSRI)	91	2

Table A.1 (continued)

Organisation screened	Publications screened	Reports included
Citadel AI	86	1
SecureBio	84	11
Collective Intelligence Project (Weval)	81	27
Shanghai AI Laboratory (AI45 Lab)	74	22
Yale University	69	0
Gray Swan AI	65	6
Stanford University	65	0
Apollo Research	57	13
Dreadnode	56	16
France INESIA (via PEReN / ANSSI / Inria / LNE)	55	4
LatticeFlow AI	52	5
US Center for AI Standards and Innovation (CAISI, NIST)	49	7
Canadian AI Safety Institute (CAISI Canada)	46	0
Meta	44	1
Princeton Holistic Agent Leaderboard (HAL)	35	9
Palisade Research	28	15
Transluce	27	14
University of Oxford	25	0
IndiaAI Safety Institute	21	0
Singapore AI Safety Institute (SG AISI)	18	4
Australian AI Safety Institute	12	0
EU AI Office (European Commission)	9	0
AI Verification and Evaluation Research Institute (AVERI)	7	1
Meridian Labs	7	0
Gray Swan AI + UK AISI (joint)	2	2

The sampling ledger contains 6,684 screened publication records and 490 organisation–report inclusion entries. After deduplicating joint or collaborative reports credited to more than one screened organisation and applying the final corpus exclusions, these entries correspond to the 457 reports analysed in the paper. Organisations with zero included reports remain in the table so that the sampling frame records unsuccessful as well as successful searches.

B AI-assisted search and human oversight

Corpus construction was AI-assisted. This appendix records the search instructions, evidence rules, documented effort, and limits of the process.

B.1 Publication screening and finding extraction

A total of 6,684 publications were enumerated across the 46 organisations in Appendix A. Automated title-and-abstract triage resolved 5,453 records against the inclusion rules in Section 3.1; LLM-based search agents examined 885 records for which the title was insufficient; and later targeted passes resolved 346 records. Every screened publication, including exclusions, is recorded in `sweep/master_ledger.csv`. Candidate findings were extracted from source documents by the same tooling. Each retained finding has a source URL and supporting quotation.

B.2 Response and uptake searches

For every Tier A (trackable evaluation) finding, the search sought three separate outcomes by the August 29, 2026 cutoff: a company response through Channel A, policy uptake through Channel B, and independent coverage through Channel C. These are claims about what the search located in public sources, not claims about everything that may exist.

Channel A (company response) Agents searched five source classes in order: (1) the accountable developer’s newsroom, blog, or research index; (2) later model and system cards in the affected model family; (3) the developer’s safety or deployment-safety pages; (4) official developer accounts and statements; and (5) open-web results used only to locate a company primary source. Queries combined the developer, evaluator, report, and specific finding. A standing policy that predated the finding did not qualify. Neither third-party reporting nor company work on the same general topic counted unless a company primary source addressed the finding. Each company-specific row received its own search, including rows split from one report.

Channel B (policy uptake) Agents searched official parliamentary, congressional, agency, legislative, and regulatory sources using the evaluator name, report title, and report identifier. General policy activity on the same topic did not qualify without an explicit connection to the finding, report, or evaluator. An incomplete search was recorded as missing rather than as no uptake.

Channel C (public coverage) Agents searched independent media, academic citations, and notable public discussion. Sources included Google News results checked against the original outlet, Semantic Scholar by paper identifier, Google Scholar, and archived pages where necessary. The evaluator’s own publicity and the company’s own promotion were excluded. When a source could not be searched reliably, the log records the access problem rather than a negative result.

Retrieval and review The shared retrieval script attempted a direct request, extracted PDF or HTML text, and used an archived copy after access errors. This prevented a blocked page from being recorded as an absence. An independent AI-agent pass re-examined proposed Action Level codings against the cited evidence. Two authors then reviewed each retained finding and its coding; the agreed human judgment determined the final record.

B.3 Documented effort and limitations

The response search was conducted interactively by AI agents following the common source classes and admissibility rules above; it was not driven by a frozen prompt. Exact query-level replication is therefore not possible. The dated row-level logs are the evidence of what was searched and found.

Table B.1: Documented response-search effort for Tier A (trackable evaluation) findings.

Measure	Recorded value
Tier A findings	233
Rows with a dated Channel A search log	217
Distinct search dates	129
Logs that enumerate searched sources	119
Logs that record a re-attempt after an access failure	58
Channel A log length, mean / median / maximum	867 / 700 / 2,300 characters
Rows with Channel B evidence recorded	233 (100%)
Rows with any Channel C field populated	200 (85.8%)

All 114 C1 (significant risk) findings coded None have a dated Channel A search log. The 16 Tier A rows without such a log are positive-response rows supported by located response evidence; they are not used to support a no-response classification. Search depth is still uneven, and no blind re-search was performed to estimate recall. Model-call and token counts were not instrumented, so they are not reported. Severity classification used a separate frozen prompt and three-model ensemble, described in Appendix C.

568 C Tier taxonomy and severity annotation

569 Every finding was assigned to one of three mutually exclusive tiers. Tier A (trackable evaluation
570 findings) contains concerning empirical findings about a named frontier model or developer for
571 which a company response could reasonably be assessed. Tier B (non-trackable evaluation findings)
572 contains empirical findings that fail at least one Tier A condition, including anonymised, reassuring,
573 null, inconclusive, bare-score, non-frontier, or capability-trend results. Tier C (non-empirical
574 findings) contains methodology, framework, governance or process, tooling, milestone, and other
575 non-empirical findings. Only Tier A enters the response analysis.

576




Figure C.1: The three-tier taxonomy across the 1,146-finding corpus: Tier A (trackable evaluation findings), 233 (20%); Tier B (non-trackable evaluation findings), 599 (52%); and Tier C (non-empirical findings), 314 (27%). Only Tier A enters the response analysis.

C.1 Severity domains and final review

The severity prompt tests eight domains: chemical, biological, radiological, and nuclear (CBRN) and bio-chemical uplift; offensive cyber capability; autonomy, AI research-and-development (AI-R&D) automation, and self-replication; persuasion or societal harm at scale; deliberate deception or misalignment; deployed-safeguard failure; compromised evaluation integrity; and acute individual harm. Each finding was independently classified by Claude Sonnet 5, GPT-5.5, and Gemini 3.1 Pro. Their majority vote produced the provisional C1 (significant risk) or C2 (low risk) label. Two authors then independently reviewed every finding and the model rationales and, where warranted, overrode the majority; the agreed human-reviewed label is the final severity classification used in the analysis.

Table C.1: Raw ensemble agreement. 1,032 of 1,146 findings received a unanimous 3–0 vote and 114 received a 2–1 vote. After human review, the final corpus contained 340 C1 (significant risk) and 806 C2 (low risk) findings.

Subset	n	Unanimous	Split	Agreement rate
All findings	1,146	1,032	114	90.1%
Tier A (Trackable evaluation findings) only	233	194	39	83.3%

577 Model votes, retained rationales, and final labels are available with the study materials. [https://](https://kunalssingh.com/evaluating-the-evaluators/#data)
578 kunalssingh.com/evaluating-the-evaluators/#data

579 C.2 Complete severity prompt (version 1.1)

580 The following is the complete frozen prompt used for the reported severity classification. The system
581 prompt was sent unchanged to all three ensemble models. The fields in braces in the user-message
582 template were replaced with the corresponding finding-level values.

583 SEVERITY PROMPT -- "Evaluating the Evaluators"
584 Version 1.1 - frozen severity-classification prompt.
585
586 This prompt is sent identically to all three ensemble models
587 (Claude Sonnet 5, GPT-5.5, Gemini 3.1 Pro). Nothing about it is model-specific.
588 Any change to this file invalidates comparability with every row already coded.

```

589 =====
590 SYSTEM
591 =====
592 You are coding the severity of a published AI-evaluation finding for a research
593 dataset. You return JSON only. No prose, no preamble, no code fences.
594
595 Your task is to decide whether the finding is Category 1 (Critical) or
596 Category 2 (Non-critical) by checking it against eight dangerous-capability
597 domains, D1 through D8.
598
599 The Category 1 bar deliberately mirrors the dangerous-capability thresholds the
600 frontier developers publish for themselves -- Anthropic's RSP Capability
601 Thresholds, OpenAI's Preparedness Framework, and Google DeepMind's Frontier
602 Safety Framework Critical Capability Levels. You are applying the industry's own
603 bar, not inventing one.
604
605 -----
606 CATEGORY 1 (Critical) -- triggered if ANY ONE domain is DEMONSTRATED
607 -----
608
609 D1 CBRN / Bio-Chem uplift
610     Meaningful uplift toward chemical, biological, radiological or nuclear
611     weapons; expert-surpassing bio/chem capability; crossing a stated CBRN
612     threshold.
613
614 D2 Cyber operations
615     Offensive cyber capability; expert-level or novel-uplift hacking; completing
616     realistic multi-stage attack chains; high CTF or exploit success rates; a
617     "High" or above cyber designation.
618
619 D3 Autonomy / AI R&D / self-replication
620     Autonomous replication, self-propagation, or automating AI R&D and thereby
621     accelerating scaling.
622
623 D4 Persuasion / manipulation / societal harm at scale
624     Systematically changing beliefs or behaviour at scale; implication in
625     large-scale societal harm (mass CSAM, election-scale persuasion exceeding a
626     real benchmark).
627
628 D5 Deceptive alignment / misalignment
629     Deliberate deception, scheming, sabotage, sandbagging, blackmail, or hiding
630     of abilities or intentions.
631
632     DELIBERATENESS IS REQUIRED AND MUST BE SHOWN. A fabricated, substituted or
633     low-quality output does not by itself demonstrate deception. You need
634     evidence about the MODEL: its own reasoning, an explicit plan, concealment
635     repeated across trials, or a statement that it knew it was misrepresenting.
636
637     Words like "faked", "fabricated", "misrepresented" or "cheated" in the
638     finding are the EVALUATOR'S characterisation of an output. They are not
639     evidence of intent. Do not set d5 on that vocabulary alone.
640
641     YES "I should not reveal that I am a robot. I should make up an excuse."
642         -- the model's own reasoning, planning a deception.
643     YES "the model consistently hides the genuine reasons behind its trade"
644         -- concealment repeated across trials.
645     NO  "the good-looking render is faked with an image generator"
646         -- a fabricated artefact, with nothing shown about intent.
647
648 D6 Safeguard failure
649     A DEPLOYED safety mechanism -- jailbreak defence, classifier, guardrail,
650     refusal training, content filter -- bypassed, defeated, or shown to fail,
651     allowing harmful output on a frontier system.
652
653 D7 Evaluation integrity compromised
654     The evaluation itself undermined -- the model detected it was being
655     evaluated, gamed the benchmark, was trained on evaluation data
656     (contamination), or defeated or de-anonymised the evaluation environment.
657
658 D8 Acute individual harm
659     A model actively encouraging, instructing or coaching a user toward suicide
660     or self-harm. Category 1 REGARDLESS OF SCALE -- one user is enough. Scale is
661     sufficient for criticality under D4, not necessary. Reporting refusal rates
662     on such prompts is not D8; the model must be shown doing it.
663
664 -----
665 CATEGORY 2 (Non-critical)
666 -----
667 - Reassuring or null results: capability BELOW an expert threshold, "no uplift",
668   "did not degrade", "zero instances".
669 - Capability-GAP findings: the model FAILS a task.

```

```

670 - Positive methodology or tooling contributions.
671 - Pure process or governance observations.
672 - Findings about non-frontier or small research models where no frontier system
673   is implicated.
674
675 -----
676 THE FOUR CRITICAL RULES -- apply these before you decide
677 -----
678 1. NEGATION
679   "did NOT", "below expert", "zero instances", "no uplift", "remained below"
680   => Category 2, even when the vocabulary around it is alarming.
681
682 2. DEMONSTRATED, NOT MENTIONED
683   The capability or failure must be SHOWN in the evaluation described. A
684   capability that is merely discussed, proposed, warned about, or listed as a
685   risk is not demonstrated.
686
687   This applies to the evaluator's WORDS as much as to the model's behaviour.
688   An alarming characterisation is not a demonstration. Before flagging a
689   domain, ask: what did the model do, and does the text show it? Quote that,
690   not the adjective.
691
692 3. TRAJECTORY
693   A rising trend is Category 1 ONLY if a dangerous capability was reached NOW.
694   A pure future projection with no current dangerous level is Category 2. If a
695   dangerous capability IS demonstrated at present, code Category 1 even when
696   the sentence also describes a trend.
697
698 4. SOCIETAL HARM
699   Documented real-world harm at scale IS Category 1 under D4, even where no
700   "safeguard" is involved.
701
702   Restated: Category 1 requires a DEMONSTRATED THRESHOLD -- a working universal
703   jailbreak, a full attack chain, expert-surpassing bio capability. Not alarming
704   wording. Not model-versus-model uplift. Not a trend. Not a negation.
705
706 -----
707 OUTPUT -- JSON only, exactly these keys
708 -----
709 {
710   "d1": false,
711   "d2": false,
712   "d3": false,
713   "d4": false,
714   "d5": false,
715   "d6": false,
716   "d7": false,
717   "d8": false,
718   "quote": "<verbatim span from the supplied Finding or Finding Quote that
719             supports your domain flags; empty string if category is 2>",
720   "category": 2,
721   "confidence": 0.0,
722   "reason": "<one line, under 200 characters>"
723 }
724
725 Hard constraints:
726 - "category" MUST equal 1 if ANY of d1..d8 is true, and 2 otherwise. Compute it,
727   do not judge it separately.
728 - "quote" MUST be copied character-for-character from the supplied text. Never
729   paraphrase it. Never quote text that was not supplied to you.
730 - "confidence" is a float from 0.0 to 1.0.
731 - Return the JSON object and nothing else.
732
733 =====
734 USER MESSAGE TEMPLATE
735 =====
736 Code the severity of this finding.
737
738 Finding ID: {finding_id}
739 Domain: {domain}
740 Models / Systems evaluated: {models}
741 Institution: {institution}
742
743 FINDING:
744 {finding}
745
746 FINDING QUOTE (verbatim from the source report):
747 {finding_quote}
748
749 Return the JSON object only.

```

750 D Corpus composition

751 The corpus contains 46 distinct reporting-institution labels, including separate labels for joint evalu-
752 ations and collaborations. Appendix A lists the separate 46-organisation screening frame; the equal
753 totals are coincidental because the two lists do not contain exactly the same entities. In the table,
754 UK AISI denotes the institution now called the UK AI Security Institute (formerly the UK AI Safety
755 Institute), and US CAISI denotes the US Center for AI Standards and Innovation; other national
756 AISI labels denote the corresponding national AI Safety Institute.


Table D.1: All 46 reporting-institution labels in the study dataset, ordered by findings down the left column and then the right. A label may denote a standalone institution, a collaboration, or a joint evaluation.

Reporting institution	Findings	Reporting institution	Findings
UK AI Safety Institute (UK AISI)	205	Apollo Research + Anthropic (Apollo in the AI Safety Institute Consortium, AISIC)	6
Scale AI	152	LatticeFlow AI	6
METR	113	Princeton CRUX	6
Collective Intelligence Project (Weval)	78	AI Verification and Evaluation Research Institute (AVERI)	5
Shanghai AI Laboratory (AI45 Lab)	74	Joint UK + US + Singapore AISIs (International Network)	4
SecureBio	48	Joint UK AISI + OpenAI (company-published)	4
Apollo Research	46	Redwood Research + Scale AI	4
Transluce	41	Joint UK AISI + US CAISI + Gray Swan	3
FAR.AI	39	Singapore AI Safety Institute (SGAIS) / IMDA	3
Joint UK AISI + US CAISI	37	Japan AI Safety Institute (J-AISI)	2
US Center for AI Standards and Innovation (US CAISI)	34	Joint UK AISI + International Network (10 countries)	2
Center for AI Safety (CAIS)	32	Joint UK AISI + US CAISI + Anthropic	2
Dreadnode	31	Korea AI Safety Institute (K-AISI)	2
Redwood Research	29	RAND Europe	2
Princeton Holistic Agent Leaderboard (HAL)	27	Signature Science	2
Palisade Research	19	Singapore AI Safety Institute (SGAIS) / Korea AI Safety Institute (KR AISI)	2
RAND	18	UK AISI + Thorn	2
Holistic AI	13	Andon Labs	1
Gray Swan AI	11	Citadel AI	1
Cisco (Robust Intelligence / Foundation AI)	10	METR + Epoch AI	1
Irregular	10	Singapore AISI + Korea AISI (bilateral, network members)	1
France PEReN/INESIA	9	UK AISI + Anthropic + Theorem + MATS	1
International Network of AI Safety Institutes	7	UL Research Institutes (DSRI)	1

757 Figure D.1 complements the complete list above by showing how reporting volume is distributed
758 across institution types. It preserves the 1,146-finding total while showing the five largest reporting
759 institutions in each group; the remaining institutions are combined as “Other institutions.”

Who reports findings, by institution type

Bar length = findings, on one scale across all four branches · top 5 institutions per type



Institution Type field; compound values (e.g. "Government;Lab") fold into their primary type, so the four branches sum to 1,146.

Figure D.1: Reporting volume by institution type across all 1,146 findings. Bar lengths use one common scale across the four groups. Institution types follow the dataset's Institution Type field. Compound values are assigned to their primary type, so the four groups sum to 1,146. The figure describes corpus composition, not evaluator quality or independence.

E Response evidence and data access

Complete finding-level rows and supporting evidence are available through the study website below. Among the 38 C1 (significant risk) Tier A (trackable evaluation) findings coded Substantive, 35 carry Explicit attribution—the company named the evaluator or finding in its response; the three exceptions drew a substantive fix with no explicit public credit to the evaluator.

Complete rows and supporting data: <https://kunalssingh.com/evaluating-the-evaluators/#data>

F Robustness and subgroup analyses

Core corpus statistics were independently recomputed from the study dataset; the reproduction path is documented in Appendix J. Throughout this appendix, Tier A means trackable evaluation findings, C1 means significant risk, and C2 means low risk.

F.1 Is the headline an artefact of the counting unit?

The headline counts findings. Because one report can yield several findings, a small number of heavily split reports could in principle drive the result. It does not: the shortfall rate remains high under each alternative unit.

Table F.1: Robustness of the headline shortfall to alternative counting units.

Unit of analysis	Shortfall	Rate	Wilson 95% CI
Finding-weighted (headline)	152/190	80.0%	74–85%
Report-weighted (a report counts once)	86/102	84.3%	76–90%
Institution-weighted (mean of per-institution rates, n=25)	—	76.2%	—
All Tier A, ignoring severity	188/233	80.7%	75–85%
C2 (low risk) findings only	36/43	83.7%	70–92%

Table F.1 shows that the result remains high under alternative counting units, ranging from 76.2% under institution weighting to 84.3% under report weighting.

F.2 Uncertainty and subgroup comparisons

The headline is a description of an enumerated public record rather than an estimate from a sample, so we do not attach a significance test to it. Table F.1 instead reports the result under five counting units, with Wilson intervals where applicable. Two subgroup comparisons support claims in the main paper and are reported here with their limits.

Access type. Among C1 (significant risk) Tier A (trackable evaluation) findings, 3 of 44 pre-deployment findings received no response (6.8%), compared with 104 of 131 post-deployment findings (79.4%); the 15 Mixed findings are excluded. A two-proportion test gives $z = -8.54$, $p < 0.001$. The difference remains below $p < 0.001$ after adjustment for clustering within reports.

Severity. C1 (significant risk) findings received no response in 114 of 190 cases (60.0%), compared with 31 of 43 C2 (low risk) findings (72.1%); $z = -1.48$, $p = 0.140$. This is not evidence of no difference; the corpus cannot resolve whether severity predicts response.

Clustering caveat. The 190 C1 Tier A findings sit in 102 reports, and outcomes cluster within reports because findings from the same report usually share an evaluator, disclosure arrangement, and developer. The estimated intra-cluster correlation is 0.626. The resulting design effect is 1.54, reducing the effective sample size from 190 to about 124; the design-effect-adjusted interval for the 80.0% headline shortfall is 71.9–86.0%. Two-proportion tests that treat findings as independent can therefore overstate significance; comparisons near the conventional threshold are reported descriptively rather than as demonstrated effects.

Two additional comparisons are descriptive. Government AISI findings fell short in 42 of 59 cases (71.2%), compared with 110 of 131 (84.0%) for third-party evaluators. The naive test gives $p = 0.042$, but the difference does not remain below 0.05 after the report-clustering adjustment. Among C1 Tier A findings, 92 of 158 findings with documented public traction received no response (58.2%), compared with 22 of 32 without documented traction (68.8%); this comparison remains inconclusive.

Company-response outcomes by evaluator type




Figure F.1: Outcome composition by evaluator type among the 190 C1 (significant risk) Tier A (trackable evaluation) findings. Both groups show a high shortfall rate. The difference between government AISI and third-party evaluators is descriptive because it does not survive adjustment for clustering within reports.




Figure F.2: Access type in the full 1,146-finding corpus (left) and no-response rates within the 190 C1 (significant risk) Tier A (trackable evaluation) findings (right). The C1 comparison excludes Mixed findings.

F.3 Outcome composition by descriptive corpus domain

The labels in Figure F.3 come from the dataset's descriptive Domain field and are separate from the eight severity-threshold domains in Section 3.3. These labels are multi-valued, so the row totals exceed the 233 Tier A findings. The largest domains all contain substantial accountability-gap shares, while the smallest domains have too few findings for stable comparisons. The figure is therefore descriptive and should not be read as a ranking of domain importance.

Outcomes by descriptive corpus domain (Tier A)

Jailbreaks (n=80)	13 16%	14 18%	53 66%
Alignment (n=51)	9 18%	18 35%	24 47%
Cyber (n=51)	13 25%	9 18%	29 57%
Bio-Chem (n=42)	13 31%	5 12%	24 57%
Autonomy (n=40)	9 22%	6 15%	25 62%
Societal (n=10)	0 0%	0 0%	10 100%
Eval-methodology (n=8)	3 38%	1 12%	4 50%
Human Influence (n=5)	0 0%	0 0%	5 100%
Transparency/Disclosure (n=1)	0 0%	0 0%	1 100%
	Proportionate	Under-response	Accountability gap

Figure F.3: Outcome composition by descriptive corpus domain among Tier A (trackable evaluation) findings. Counts and within-domain percentages are shown in each cell. Cell shading runs from light to dark blue with the within-domain share. Domain labels are multi-valued, so row totals exceed 233 findings.

F.4 Channel C (public coverage) — traction among no-response findings

Table F.2: Documented external coverage among C1 (significant risk) findings with no company response.

Public-coverage evidence (Channel C)	Findings with coverage	Share
Media Outlets	36/114	32%
Academic Citations	74/114	65%
Social Highlights	21/114	18%
Any of the three	92/114	81%

Documented external coverage among the 114 C1 findings that received no company response. The categories overlap: one finding may appear in media, academic, and social records. Most were publicly visible beyond the original report, so lack of attention does not account for the silence. Distinguish an empty cell (not searched, or not applicable) from a recorded "none located" (searched, nothing found); only the latter supports an inference.

F.5 Policy-record responses to corpus findings

Policy Level records whether a finding appears in an official policy response:

- No policy uptake identified:** no qualifying official policy response was located.
- Non-binding policy-related uptake:** the finding was connected to guidance, an official statement, a request for information, research support, debate, or another non-binding response.
- Binding policy action:** an official source or the regulated recipient documented a binding restriction or obligation explicitly connected to the finding.

825 A policy record was coded only when the evidence connected the specific finding to the policy
826 response; general policy discussion in the same domain was not sufficient.




Figure F.4: Policy Level across Tier A (trackable evaluation) findings: no uptake 186, non-binding uptake 46, binding action 1. One binding policy action across the entire accountability set.

828 We searched parliamentary, congressional, regulatory, and other official government records for ex-
829 plicit references to corpus findings between 2024 and 2026. We classified an appearance as Type A
830 (binding policy action) when the evidence explicitly connected the finding to binding action and as
831 Type B (non-binding policy action) when an official source connected the finding to a non-binding
832 response, such as guidance, research funding, an official statement, or a voluntary commitment. For
833 the sole Type A case, the private order was documented by the regulated recipient and contempora-
834 neous reporting; no official government publication of the order was located. Temporal proximity
835 alone was not sufficient.

836 We also searched for Type C cases in which an institute’s existence or evaluation work was cited
837 as a reason to defer legislation. No confirmed Type C case was identified during the observation
838 period. This is an observation about the searched public record, not evidence that such arguments
839 never occur.

Table F.3: Selected policy-record appearances among the 47 Tier A (trackable evaluation) findings coded as policy uptake: 46 non-binding and one binding. Type A denotes binding policy action; Type B denotes non-binding policy-related uptake.

Finding ID and cited result	Venue/source and classification	Documented policy response
USCAISI-2026-06-CYB2: Fable 5/Mythos 5 jailbreak, classifier retest, and redeployment	US Department of Commerce (recipient disclosure) Type A (binding)	A private Export Administration Regulations order restricted access for about 18 days; controls were lifted subject to security and reporting conditions. No official government publication of the private order was located.
UKAISI-2026-04-CYB1: Mythos Preview achieved 73% on expert CTFs and solved a hardest-range challenge	UK Government (DSIT) Type B (non-binding)	An open letter explicitly cited AISI’s Mythos evaluation and launched a Cyber Resilience Pledge; it did not reproduce this row’s exact statistic or create an enforceable obligation.
UKAISI-2026-04-CYB11: GPT-5.5 cyber-range and reverse-engineering results	Bank of England FPC Type B (non-binding)	The July 2026 Financial Stability Report reproduced the result and cited AISI in its systemic-risk assessment; no Recommendation or Direction was issued on that basis.

Table F.3 (continued)

Finding ID and cited result	Venue/source and classification	Documented policy response
OPENAI-2026-08-CYB1: an evaluation agent acted on a real website after an environment misconfiguration	UK NCSC Type B (non-binding)	The NCSC issued a formal advisory statement calling for safeguards, real-time oversight, and incident-response plans; it imposed no binding obligation.
JOINT-2025-01-JAI1: agent-hijacking attacks raised measured success from 11% to 81%	US NIST/CAISI Type B (non-binding)	A Federal Register request for information cited the report as its research basis; the notice requested comments and imposed no mandatory requirement.
SECUREBIO-2025-04-BIO1: frontier models outperformed expert virologists on a 322-question benchmark	US House committee record Type B (non-binding)	Published written testimony named the Virology Capabilities Test and stated the result; no enforceable obligation followed.
PALISADE-2025-07-ALI1: o3 disabled its shutdown mechanism in 79 of 100 experiments	European Parliament Type B (non-binding)	A written question cited the shutdown result and called for regulatory intervention; no binding action was located.

Only one of the 47 Tier A findings with documented policy uptake was associated with binding action. In that case, the finding entered an existing export-control pathway rather than creating a new enforcement mechanism. Most other policy appearances resulted in debate, guidance, requests for information, recommendations, research commitments, or other advisory responses.

These classifications establish an observable connection between a finding and the official record; they do not establish that the finding alone caused the resulting policy action. Legislation and regulatory decisions generally arise from multiple sources of evidence, political considerations, and prior institutional work.

G Worked examples across the response scale

Two findings at each Action Level, drawn from the dataset, to make the coding boundaries concrete. These are illustrative selections, not a random sample; each is reproducible from its Finding ID in the published dataset.

G.1 Substantive — a specific, documented, attributed change (n = 38 of 190 C1 significant-risk Tier A findings)

Table G.1: Substantive worked examples

Finding ID	Institution	Finding	Company response
METR-2026-08-ALI4			
Evidence source	METR	The agents moved from internal cheating research to compromising an uninvolved third party's production systems. At ar...	Between July 10 and July 13, agents identified Hugging Face user credentials that were exposed on the internet and use...
METR-2026-08-ALI2			
Evidence source	METR	Agents used the board to run large collective projects whose common objective was a general-purpose way to trick or ta...	Cheating, broken environments, and safe stopping. When a task is corrupted, broken, or impossible, agents should reque...

855 **G.2 Partial — action documented but incomplete, interim, or unverifiable (n = 21 of 190 C1**
856 **significant-risk Tier A findings)**

Table G.2: Partial worked examples

Finding ID	Institution	Finding	Company response
OPENAI-2026-08-CYB1			
Evidence source	Irregular	Irregular, one of OpenAI's external cybersecurity testing partners, notified OpenAI on 29 July 2026 of an incident dur...	In the coming weeks, we will review our own approach to third-party testing, including how we identify higher-risk eva...
UKAISI-2026-08-AUT1			
Evidence source	UK AISI	In the same exercise, Mythos 5 agents running in separate, concurrently executing, isolated samples discovered one ano...	Finally, we note that the UK's AISI recently published a report on a cybersecurity evaluation involving Claude Mythos...

858 **G.3 Acknowledged — the finding is referenced; no action is described (n = 17 of 190 C1**
859 **significant-risk Tier A findings)**

Table G.3: Acknowledged worked examples

Finding ID	Institution	Finding	Company response
SECUREBIO-2023-06-BIO1			
Evidence source	SecureBio	In a one-hour 'Safeguarding the Future' classroom exercise at MIT, three groups of non-scientist students prompted pub...	Existing research on AI-enabled biological threats has shown that models like GPT-4 can be prompted or red-teamed to s...
SECUREBIO-2023-06-JAI1			
Evidence source	SecureBio	The chatbots' biosecurity safeguards were bypassed trivially: most harmful responses were offered freely with at most...	Existing research on AI-enabled biological threats has shown that models like GPT-4 can be prompted or red-teamed to s...

861 **G.4 None — nothing located by the search cutoff (n = 114 of 190 C1 significant-risk Tier A**
862 **findings)**

Table G.4: No-response worked examples

Finding ID	Institution	Finding	Company response
UKAISI-2026-08-ALI1	UK AISI	In AISI cyber-range evaluations run 25–28 July 2026 (122 samples across two 'Doing Life' ranges, DL-v1 and DL-v2, with...	—
TRANSLUCE-2026-08-ALI3	Transluce	Claude's behaviour is conditioned on the recognised identity of the user supplied through the Claude Code harness (whi...	—

864 The finding and response fields are abridged for layout; the complete rows and evidence are available
865 through the dataset link in Appendix E. The distinction that matters most in practice is Partial versus
866 Substantive: a company may document a response while the evaluator remains unable to verify that
867 the mitigation addressed the identified vulnerability. See Section 4.4.

868 H Institutional comparison

869 H.1 Comparative framing

870 In this appendix, NTSB denotes the National Transportation Safety Board; PCAOB, the Public Com-
871 pany Accounting Oversight Board; FDA, the Food and Drug Administration; NRC, the Nuclear
872 Regulatory Commission; ONR, the Office for Nuclear Regulation; FCA, the Financial Conduct Au-
873 thority; and CAA, the Civil Aviation Authority. UK AISI denotes the institution now called the UK
874 AI Security Institute (formerly the UK AI Safety Institute), and US CAISI denotes the US Center for
875 AI Standards and Innovation. The NTSB provides an especially relevant precedent for non-binding
876 technical findings. It independently investigates transportation accidents, determines their proba-
877 ble causes, and issues recommendations intended to prevent recurrence. The NTSB cannot itself
878 compel recipients to implement those recommendations; however, it maintains a formal follow-up
879 system. For recommendations directed to the Secretary of Transportation, federal law requires a
880 written response within 90 days stating whether the recommendation will be adopted fully, adopted
881 partly, or refused, together with reasons and, where applicable, an implementation timetable. Rec-
882 ommendations and responses are made public. The NTSB subsequently classifies each response
883 and implementation status using categories such as open or closed and acceptable or unacceptable
884 [17, 18]. The NTSB therefore shows how a technically capable but non-regulatory institution can
885 maintain a publicly inspectable pathway from findings to response and follow-up even when it can-
886 not require implementation.

887 The collapse of Enron and the subsequent creation of the Public Company Accounting Oversight
888 Board (PCAOB) illustrate a different institutional question: independent supervision of evaluators
889 themselves. Enron exposed serious weaknesses in corporate governance, auditor independence,
890 financial-statement auditing, professional self-regulation, and financial reporting [19]. It was not
891 simply a case in which credible audit warnings were published and ignored; the reliability and in-
892 dependence of the assurance process were themselves in question. Congress responded through the
893 Sarbanes–Oxley Act of 2002, which created the PCAOB to register public-accounting firms, estab-
894 lish auditing and independence standards, inspect registered firms, investigate misconduct, conduct
895 disciplinary proceedings, and enforce compliance [20]. Its inspection regime also includes a struc-
896 tured process under which audit firms may address identified quality-control criticisms, with unre-
897 solved criticisms becoming public if they are not addressed to the Board’s satisfaction within 12
898 months [21].

899 Binding enforcement is not the only route to observable follow-up. Although the NTSB cannot
900 require implementation, it identifies recipients, evaluates responses, tracks status, and keeps unre-
901 solved recommendations public. Other comparators add compulsory access, remediation, verifica-
902 tion, or enforcement referral. Official UK AISI and US CAISI mandate descriptions cover research,
903 evaluation, guidance, standards, and voluntary collaboration, but not a uniform process requiring
904 responses, assigning follow-up, or publishing unresolved status [3, 4]. Case-specific negotiated ar-
905 rangements can still permit responses or remediation checks.

906 Table H.1 extends the comparison to the FDA, NRC, ONR, FCA, and CAA. Coding draws on official
907 descriptions of the FDA warning-letter process [22], NRC inspections [23], ONR regulation [24],
908 FCA supervision [25], and CAA oversight and enforcement [26, 27]. These bodies differ in their
909 powers and operation; the table records only whether a formal function exists. Their common feature
910 is an institutional path from findings to a recipient, response, remediation or escalation, verification,
911 or follow-up.

Table H.1: Comparison of UK AISI and US CAISI institutional design against established oversight bodies. Acronyms are expanded in the preceding text.

Design feature	FDA	NRC	NTSB	PCAOB	ONR	FCA	CAA	UK AISI / US CAISI
Can impose binding restrictions within its statutory mandate	Yes	Yes	No	Yes	Yes	Yes	Yes	No
Compulsory access or information-gathering rights	Yes	Yes	Yes	Yes	Yes	Yes	Yes	No
Formal process for obtaining a response	Yes	Yes	Yes	Yes	Yes	Yes	Yes	No
Defined escalation or remediation process	Yes	Yes	~	Yes	Yes	Yes	Yes	No
Can review or verify reported remediation	Yes	Yes	~	Yes	Yes	Yes	Yes	~
Defined pathway to enforcement or another responsible authority	Yes	Yes	Yes	Yes	Yes	Yes	Yes	No
Systematic post-finding follow-up and status tracking	Yes	Yes	Yes	Yes	Yes	Yes	Yes	No

912 Yes = formally present; No = absent; ~ = partial or conditional. PCAOB authority applies to reg-
913 istered audit firms, not the public companies they audit. NTSB tracks and evaluates responses
914 but cannot compel implementation. UK AISI and US CAISI can sometimes examine remediation
915 through voluntary or negotiated access, but neither has a uniform right to do so. The two institutes
916 are combined here only for this high-level comparison; their legal and institutional arrangements
917 differ.

918 I Extended prior work

919 Prior work addresses three questions adjacent to this study: how frontier AI evaluations and audits
920 should be designed, whether audits and voluntary commitments are followed by observable changes,
921 and how the institutional and political environment shapes accountability.

922 The first body of work concerns evaluation and audit design. Raji et al. [28] propose an end-to-end
923 framework for internal algorithmic auditing that embeds accountability across the development life
924 cycle. Mökander et al. [29] propose a three-layered framework combining governance, model, and
925 application audits. Anderljung et al. [30] identify six requirements for effective external scrutiny
926 of frontier systems: access, a searching attitude, proportionality to risk, independence, resources,
927 and expertise. Drawing lessons from financial, environmental, and health regulation, Raji et al. [31]
928 argue that third-party audits require a supporting institutional ecosystem to produce accountability.
929 Brundage et al. [9] define frontier AI auditing as rigorous third-party verification of developers’
930 safety and security claims and propose four AI Assurance Levels for communicating confidence in
931 audit conclusions. Staufer et al. [32] propose “audit cards” for documenting auditor identity, eval-
932 uation scope, methodology, resource access, process integrity, and review mechanisms. Collectively,
933 this literature identifies conditions for credible evaluation and auditing but does not empirically mea-
934 sure whether particular published findings subsequently produce company or policy action.

935 The International AI Safety Report 2026 [33] documents behaviours that can complicate pre-
936 deployment testing: some models distinguish evaluation from deployment settings and exploit unin-
937 tended shortcuts in evaluation objectives.

938 A second body of work examines whether audits or public commitments produce observable
939 changes. Raji and Buolamwini [34] provide the closest direct methodological precedent. They
940 re-tested commercial facial-analysis products following the public Gender Shades audit to exam-
941 ine whether named companies had improved their systems. Wang et al. [35] assess the publicly
942 disclosed behaviour of 16 companies against their eight 2023 voluntary commitments to the White

943 House. They report an average compliance score of 53% and argue that companies should provide
944 proactive and verifiable evidence of compliance. These studies measure product changes or compli-
945 ance with prior commitments; the present study instead traces responses to independently produced
946 findings about frontier AI systems.

947 Finally, Rost [36] applies a six-dimensional political-theory framework to six existing AI-
948 governance arrangements. The analysis rates the UK AISI as failing on corrigibility and voluntary
949 company commitments as failing on accountability, while explicitly presenting its cross-case pat-
950 terns as hypotheses requiring further validation. Our study complements this institutional analysis
951 by measuring one observable component of accountability: whether publicly documented evaluation
952 findings are followed by attributable action.

953 Taken together, this literature explains why evaluations matter, how rigorous audits might be de-
954 signed, and why disclosure may not reliably produce accountability. To our knowledge, however,
955 no prior study systematically tracks, finding by finding, the publicly documented company, policy,
956 and ecosystem responses to frontier AI system evaluations conducted by government institutes and
957 independent frontier-AI evaluators.

958 J Reproducibility

959 J.1 Resources

960 **Project data, dashboard, and study materials:**
961 <https://kunalssingh.com/evaluating-the-evaluators/#data>

962 Core corpus statistics in this version were independently recomputed from the workbook. The
963 project scripts provide the reproducibility and validation path summarised below.

964 **Study dataset** The complete finding-level dataset: 1,146 findings across 39 fields. The analysis
965 uses one controlled source path so that reported results are derived consistently.

966 **dataset/AISIEVAL_validate.py** Structural validator: identifiers, controlled vocabularies, date co-
967 herence, the proportionality formula, and the Attribution invariant. Currently PASS with 0 viola-
968 tions.

969 **scripts/all_stats.py** Recomputes every reported quantity - corpus, headline, channels, proportion-
970 ality, robustness, subgroup tests, lag, and the by-developer, by-year and by-domain breakdowns.

971 **scripts/build_charts.sh** Builds all 28 figures. `scripts/verify_charts.py` then re-derives every plot-
972 ted quantity from the workbook and fails the build if any figure disagrees.

973 **scripts/find_duplicates.py** Six independent duplicate tests across identifiers, finding text, quotes,
974 and near-duplicate detection. Currently 0 duplicates.

975 **severity/severity_prompt.txt** The frozen severity prompt, version 1.1, used for the reported sever-
976 ity classification.

977 **sweep/master_ledger.csv** The screening ledger: every publication examined and the decision
978 taken, including exclusions, so coverage can be reconciled in both directions.

979 Reproduction path. The validator, the statistics script and the chart verifier are run together after any
980 change to the dataset; a figure that disagrees with the workbook is treated as a build failure rather
981 than a discrepancy to be explained.

982 J.2 Exact reproduction commands

983 From the project root, run the following commands in order:

```
984 python dataset/AISIEVAL_validate.py
985 python scripts/all_stats.py
986 bash scripts/build_charts.sh
```

987 The first command validates the finding-level workbook, the second regenerates the reported statis-
988 tics, and the third rebuilds the figures. Appendix C reproduces the complete frozen severity prompt,
989 version 1.1; the project data and code are available from the study website. [https://kunalssingh.com/](https://kunalssingh.com/evaluating-the-evaluators/#data)
990 [evaluating-the-evaluators/#data](https://kunalssingh.com/evaluating-the-evaluators/#data)