
Do third-party frontier AI evaluations matter?

Kunal Singh^{1,*}, Max Kamache¹, Jonas Raedler¹, Stephen Casper²

¹*MATS Research*

²*Harvard Kennedy School*

**Corresponding author: singhkunal9373@gmail.com*

Abstract

Government AI institutes and independent evaluators often evaluate frontier AI systems to reduce reliance on self-reporting by frontier developers. Yet identifying and reporting a risk does not itself mitigate that risk. Here, we examine whether published third-party AI evaluation findings concerning named frontier systems are acted upon. We examine publicly documented findings from evaluations conducted by government AI institutes and independent evaluation organisations, covering reports available through August 2026. We trace downstream responses through three channels: company responses and model updates; parliamentary and regulatory action; and documented media and academic coverage. The corpus contains 1,144 findings, of which 231 were concerning empirical results about named systems for which a company response could reasonably be expected. Of the 188 that met our predefined significant-risk threshold, 150 (80%) lacked a specific, publicly documented company response meeting our proportionality criteria. Together, these findings suggest that, based on publicly available information, there is only limited reason to trust that third-party evaluations of frontier AI systems consistently lead to proportionate, attributable, and verifiable responses. We compare the accountability structures currently in place for frontier AI systems with those established by oversight bodies in other industries, including food, drugs, energy, transportation, and finance. We conclude that formal mechanisms for access, response, remediation, verification, and follow-up inspired by governance in these other sectors could increase the consistency with which AI evaluations translate to meaningful action.

1 Introduction

Frontier AI system evaluations have become a prominent component of AI governance. Government AI institutes and independent evaluators test advanced systems for dangerous capabilities, safeguard failures, and other risks, often with access and technical expertise unavailable to the general public. These evaluations can provide evidence about an AI system’s risk beyond the boundaries of the developers’ own disclosures [23].

However, identifying a risk is not the same as reducing it. An evaluation can document a serious vulnerability without compelling a developers’ response, which could lead to delay in deployment, additional reviews, or verification of remediation. Therefore, the practical value of the evaluation depends partly on what happens after the finding is communicated or published.

The importance of a finding is not determined by whether it is followed by any publicly documented action, and this paper does not conflate the two. An evaluation that does not identify a concern or actionable problem is not necessarily unsuccessful. Conversely, the absence of publicly documented action does not establish that no action occurred. Some effects of evaluations may be invisible to this study. Anticipation of an evaluation may prompt a company to conduct additional due diligence

38 or address a problem before testing begins. Evaluators can sometime communicate their findings
39 privately, and companies may implement changes without publicly attributing them. These preventive,
40 private, or unattributed responses may leave no public trace and therefore would not be captured in
41 our data.

42 Finally, we compare the frontier-AI evaluation ecosystem with existing oversight institutions in
43 pharmaceutical regulation, nuclear safety, transportation, financial, and related domains. These
44 institutions differ substantially in their mandates and enforcement powers, but they employ formal
45 mechanisms that connect technical findings to follow-up: information-access rights, designated
46 recipients, response expectations or obligations, remediation pathways, verification of corrective
47 action, escalation, status tracking, and continuing review. These comparisons are not presented as
48 direct templates for AI governance, but they illustrate how evaluations can be embedded in processes
49 designed to ensure that identified shortcomings are addressed or remain visible until formally resolved.

Table 1: Outcomes among the 188 C1 (significant-risk) Tier A (trackable evaluation) findings.

Outcome	Count	% of C1	What it means
Accountability gap (no action)	112	60%	Alarming finding concerning a named company, with no publicly documented company response.
Under-response (partial / acknowledged)	38	20%	The company acknowledged the finding or described a response, but no complete and verifiable mitigation was documented.
Proportionate	38	20%	A specific, documented response meeting the proportionality criteria; attribution is assessed separately.
Gap total (rows 1 + 2)	150	80%	Either no response or an under-response; no documented company response met the proportionality standard.

51 2 Background

52 2.1 Government AI Institutes

53 The first wave of government AI institutes, established in the UK, the United States, and Japan in 2023
54 and 2024, consisted of technical government bodies with mandates focused on the safety of advanced
55 AI systems. Although they differ in their institutional arrangements, they shared a similar emphasis
56 on safety evaluations supported by three core functions: research, standards, and cooperation [2].

57 The first-wave model subsequently developed into a broader and more institutionally varied ecosystem
58 of government bodies conducting or coordinating AI safety research, evaluation, measurement,
59 and standards work. These bodies differ in mandate, structure, system access, and relationship
60 to regulators. This paper therefore refers to “government AI institutes” as an umbrella term for
61 government bodies that conduct or coordinate external evaluations of advanced AI systems, without
62 assuming that they are institutionally equivalent.

63 First-wave AI Safety Institutes (AISIs) were designed primarily as technical rather than regulatory
64 institutions. Their evaluations could inform AI developers, policymakers, and standards develop-
65 ment, but the institutes had no regulatory powers and their findings did not themselves compel any
66 remediation [2].

67 Government AI institutes are therefore one component of the external evaluation ecosystem (evalua-
68 tions conducted outside of frontier-model developers). A second component consists of independent
69 evaluators operating outside both government and frontier-model developers.

70 2.2 Independent AI system evaluators and third-party auditing

71 Independent evaluators are defined as non-government organizations (including nonprofit organi-
72 zations, for-profit evaluation firms, and academic research groups) that are institutionally separate
73 from the developer whose system they evaluate and that conduct and publicly report frontier-AI
74 evaluations. They include generalist evaluators such as METR, Apollo Research and domain-specific
75 organizations such as SecureBio. Independent evaluators also test frontier systems for dangerous

capabilities, safeguard failures, and other risks, but differ from government institutes in institutional status, funding, access arrangements, and public responsibilities. Their access may depend on voluntary developer cooperation, contractual arrangements, application programming interfaces, or publicly available models, allowing some to conduct pre-deployment evaluations while limiting others to post-deployment testing. At the same time, independent evaluators can provide specialized expertise and scrutiny outside both government and frontier-model developers. Recent work therefore treats evaluator independence, system access, methodological rigor, reporting arrangements, and audit standards as important dimensions of third-party auditing [7–10]. This paper treats government institutes and independent evaluators as distinct institutional categories while applying the same public-record response framework to eligible findings from both.

3 Methodology

3.1 Data collection

We constructed the sampling frame by identifying 46 government bodies and independent evaluators that conducted frontier-AI evaluations and publicly reported their methods and results. We screened 6,684 publications and identified 474 reports containing sufficiently documented evaluations of advanced AI systems. The earliest included report was published on November 16, 2022, and the public-record cutoff was August 29, 2026. From these reports, we coded 1,144 distinct findings. The corpus includes adverse findings about named systems, findings about anonymised systems, reassuring or null results, capability trends, and methodology or governance findings. Section 3.2 places them in three tiers. Only the 231 Tier A (trackable evaluation) findings—adverse findings about a named model or developer—enter the response and proportionality analyses. The sampling frame and the 46 reporting-institution labels, including separate labels for joint evaluations, are listed in Appendices A and D.

The unit of analysis is an individual finding rather than a report. A finding had to be asserted by the evaluator, supported by measured results, observed behaviour, or a documented process fact, and independently codeable. Opinions, recommendations, announcements, and plans were excluded. When a report showcases a substantively distinct result, for example, a baseline safeguard-failure rate and the effectiveness of a particular bypass technique, those results were coded as separate findings. We always coded results concerning different developers as separate findings. We combined results concerning multiple models from the same developer when they documented the same issue and would reasonably call for the same response, but coded them separately when they reported substantively distinct issues or would call for different responses. Each observation received a unique Finding ID, while findings drawn from the same source retained a common Report ID.

The corpus is restricted to evidence available in the public record. Findings communicated privately but not subsequently disclosed cannot be observed. We also excluded non-English-language outputs where a sufficiently reliable English version was unavailable. The resulting corpus should therefore be interpreted as a structured collection of publicly accessible, predominantly English-language evidence, not as a complete census of every evaluation performed by the included institutions.

3.1.1 Tooling and human oversight

Corpus construction was AI-assisted. Publication screening was performed by automated title-and-abstract triage over the 6,684 enumerated publications, examined by large language model (LLM) search agents against the inclusion rules of Section 3.1; candidate findings were extracted from source documents by the same tooling. Searches for Channel A (company responses), Channel B (policy uptake), and Channel C (public coverage) were likewise distributed across agent-dispatched tasks following the response-channel definitions and evidence rules described in Section 3.4. Severity classification is a three-model ensemble, described in Section 3.3.

All inclusion decisions, tier assignments and Action Level codings were made or confirmed by the authors against the coded evidence, and every finding retained in the corpus carries a source URL and a verbatim quotation recorded at extraction. Severity is the exception: it is assigned by ensemble majority vote and overridden wherever required by the authors. Appendix B documents the tooling, dates, query strategy and review procedure in full, together with the known limits of the verification performed.

3.2 Defining the Accountability Set

We assigned each finding to one of three mutually exclusive tiers based on whether it reported an adverse or safety-relevant result and permitted developer-specific response analysis.

Tier A (Trackable evaluation findings): Response-trackable adverse findings. A finding entered Tier A (Trackable evaluation findings) when it: (1) reported a specific adverse or safety-relevant result; (2) identified a named model, system, or developer rather than only an anonymised placeholder; and (3) concerned an issue within the developer’s capacity to address, such that a company response could reasonably be assessed. These 231 findings constitute the accountability set and are the only findings included in the Action Level and proportionality analyses.

Tier B (Non-trackable evaluation findings): Empirical findings that did not meet at least one Tier A condition. These include anonymised-system results, reassuring or null results, bare scores or rankings without a concerning threshold, non-frontier systems, capability trends, inconclusive results, and findings for which a company response could not reasonably be assessed.

Tier C (Non-empirical findings): Methodology, framework, governance or process, tooling, milestone, and other findings that did not report empirical model behaviour or capability.

All three tiers remain in the descriptive corpus. Only Tier A (Trackable evaluation findings) enters the company-response and proportionality analyses.

3.3 Severity of risk classification

A finding was classified as C1 (significant risk) when the reported evidence met at least one significant-risk threshold across eight risk domains: chemical, biological, radiological, and nuclear (CBRN) uplift; offensive cyber capability; autonomy, self-replication, or AI research-and-development (AI-R&D) automation; persuasion or societal harm at scale; deliberate deception or misalignment; deployed-safeguard failure; compromised evaluation integrity; or acute individual harm. Findings could be assigned to more than one domain. A finding was classified as C2 (low risk) when the reported evidence did not meet any of these significant-risk thresholds. This does not necessarily mean that the finding was unimportant or that no response was warranted. The eight domains draw on two sources: published developer frameworks and additional risk criteria developed for this study. Domains D1–D3—CBRN uplift, offensive cyber capability, and autonomy including AI-R&D automation and self-replication—correspond to dangerous-capability thresholds that frontier developers publish for their own systems in the Anthropic Responsible Scaling Policy [3], the OpenAI Preparedness Framework [4], and the Google DeepMind Frontier Safety Framework [5]. Domains D4 and D5—persuasion or societal harm at scale and demonstrated deceptive misalignment—correspond to categories that those frameworks treat less uniformly. Domains D6–D8—deployed-safeguard failure, compromised evaluation integrity, and acute individual harm—are additional criteria developed for this study. They describe failures observed in the evaluation literature that the developers’ capability thresholds do not cover. All threshold wording and boundary rules are the authors’ operationalization.

The majority vote of the three models (Claude Sonnet 5, GPT-5.5, and Gemini 3.1 Pro) provided the initial severity classification. A human annotator then reviewed every finding and its supporting evidence, including the individual model votes and rationales. Where warranted, the annotator overrode the model majority; the human-reviewed label was the final severity classification used in the analysis. Individual model votes and model-generated reasons were retained so that model disagreements remained observable. The complete annotation prompt, eight domains, boundary rules, and output format are reproduced in Appendix C.

Across the full corpus, 1,030 of 1,144 findings (90.0%) received a unanimous 3–0 classification, while 114 (10.0%) were determined by a 2–1 vote. Among the 231 Tier A (Trackable evaluation findings) findings, 192 (83.1%) received a unanimous classification and 39 (16.9%) were determined by a 2–1 vote. The final corpus contained 333 C1 (significant risk) and 811 C2 (low risk) findings. Within Tier A (Trackable evaluation findings), 188 findings were classified C1 (significant risk) and 43 were classified C2 (low risk). Agreement was high across the full corpus and within Tier A (Trackable evaluation findings).

179 3.4 Response Evidence and Outcome Coding

180 3.4.1 Evidence channels

Table 2: Evidence channels searched separately through August 29, 2026. Channels B and C do not count as a company response.

	Channel	Evidence examined	Role in the analysis
181	Channel A (company response)	Company statements, system cards, and safety or deployment pages.	Attribution and Action Level.
	Channel B (policy uptake)	Official parliamentary, legislative, regulatory, and government records.	Policy Level.
	Channel C (public coverage)	Independent media coverage, academic citations, and notable public discussion.	External visibility.

182 The resulting public-record evidence was coded into four study-defined measures: Attribution, Action
183 Level, Policy Level, and Proportionality.

184 3.4.2 Attribution

185 Attribution records whether an identified response was publicly connected to the relevant finding. (a)
186 Explicit attribution: The responding company directly referenced the finding, report, evaluation result,
187 or evaluating institution, or stated that its response was informed by that evidence. (b) No explicit
188 attribution: A qualifying response was identified, but the company did not publicly state a connection
189 to the finding. (c) No response located: No qualifying company response was identified, so attribution
190 does not arise. Attribution is assessed only where a company response exists. Cases with no identified
191 response are recorded under the third category rather than left blank, so that a searched-and-empty
192 result is distinguishable from an uncoded one. Attribution measures the company's public explanation
193 and does not independently establish causation.

194 3.4.3 Action Level

195 Action Level records the strength of the publicly documented company response identified through
196 Channel A (Company response). (a) None: No public response could be identified through the
197 search by the cutoff date. (b) Acknowledged: The company recognised the finding or underlying
198 problem but specified no action. (c) Partial: The company documented an action that addressed
199 only part of the identified problem or was expressly described as interim, limited, or incomplete.
200 (d) Substantive: The company documented a specific mitigation, model change, safeguard, access
201 restriction, or deployment decision that directly addressed the identified problem. Action Level
202 measures the content of the public response. It does not establish whether the action was successfully
203 implemented or effective in reducing the risk. Only public primary documentation issued by the
204 responding company was admitted as Channel A evidence; evaluator statements, news coverage, and
205 other third-party sources did not themselves qualify as company responses.

206 3.4.4 Policy Level

207 Policy Level records whether the finding received documented uptake through Channel B (policy
208 uptake) and whether it created an enforceable requirement. The categories are No policy uptake
209 identified, Non-binding policy-related uptake, and Binding policy action.

210 3.4.5 Proportionality

211 Proportionality combines two independently assigned classifications: finding severity and company
212 Action Level. In Figure 1, the rows show Severity (C1 significant risk or C2 low risk), and the columns
213 show the documented company-response level. The matrix shows the severity-dependent rule used
214 to classify every Tier A (Trackable evaluation findings) finding. Equivalently, a C1 (significant
215 risk) finding requires a Substantive response to meet the proportionality standard, whereas a C2
216 (low risk) finding requires at least a Partial response. Acknowledgment without action never meets
217 the standard, and the absence of a company response always produces a no-action gap. This is a
218 study-defined measure of documented response proportionality. It evaluates whether the publicly
219 described response met the minimum Action Level associated with the finding's severity classification.
220 It does not establish that the response was effective, that remediation was successfully implemented,
221 or that the underlying risk was eliminated.

SEVERITY	COMPANY RESPONSE LEVEL			
	Substantive	Partial	Acknowledged	None
C1 <i>significant risk</i>	PROPORTIONATE	UNDER-RESPONSE	UNDER-RESPONSE	ACCOUNTABILITY GAP
C2 <i>low risk</i>	PROPORTIONATE	PROPORTIONATE	UNDER-RESPONSE	ACCOUNTABILITY GAP

Figure 1: Proportionality coding rule. Rows show severity; columns show company-response level; cells show the resulting outcome. The matrix defines the rule, not observed results.

4 Results

We first report company-response outcomes among the 188 C1 (significant-risk) findings in Tier A (trackable evaluation findings), which constitute the paper’s primary accountability analysis. We then examine differences by evaluation-access type, public coverage of findings without company responses, and the extent to which reported company responses could be independently verified.

4.1 Company Responses to Significant-Risk Findings

Restricted to C1 (significant-risk) findings—the comparison relevant to the paper’s central accountability question—150 of 188 findings (80%; Wilson 95% CI 73–85%) received either no publicly documented company response or an inadequate one. This combined figure, rather than the raw 62% no-response rate across all 231 Tier A (trackable evaluation) findings, is the paper’s central result. Because it is calculated only over C1 findings, it is not mechanically affected by the number of C2 (low-risk) findings included in the corpus. Of the 76 C1 findings with a located company response, 64 had explicit attribution and 12 did not. In Channel B (policy uptake), 41 C1 findings received non-binding policy uptake, one received binding action, and 146 had no identified policy uptake. Robustness checks and subgroup tests are reported in Appendix F.

4.2 Outcomes by Evaluation Access Type

Among the Tier A (trackable evaluation) findings classified as either pre- or post-deployment, the substantive-response rate was 54% for the 54 pre-deployment findings and 6% for the 160 post-deployment findings. Pre-deployment findings were therefore roughly nine times as likely to receive a substantive response. No-response rates also differed substantially: 5.6% for pre-deployment findings and 82.5% for post-deployment findings. This difference remained significant after adjustment for clustering within reports ($p < 0.001$; Appendix F.2). These comparisons describe an association and should not be interpreted causally.

This descriptive association is consistent with pre-deployment access facilitating company action, although access type is correlated with evaluator, disclosure arrangements, domain, and other institutional differences. One plausible mechanism is visible in the case-level record: pre-deployment findings are communicated while the deployment decision window remains open and are frequently published under coordinated-disclosure arrangements. Forty-nine of the 88 recorded response lags were zero days, meaning that the finding and company response were published together. This remains a plausible interpretation rather than a mechanism established by the comparison itself.

4.3 Public Coverage of Findings Without Company Responses

Of the 112 C1 (significant risk) Tier A (trackable evaluation) findings classified as Accountability gap (no action), 89 (79%) have at least one mainstream press item or academic citation recorded. Including notable public discussion, 90 (80%) have evidence in at least one Channel C (public coverage) field. Thus, the absence of a publicly documented company response was not generally

258 confined to findings with no observable external coverage. This evidence does not establish that the
259 relevant company saw the coverage, but it shows that many no-response findings were publicly visible
260 beyond the original evaluation report. The channel-by-channel breakdown is reported in Table F.2;
261 Appendix F.2 reports the inconclusive comparison with findings lacking documented traction.

262 4.4 Independent Verification of Company Responses

263 A company response does not necessarily permit independent verification. During its pre-deployment
264 evaluation of GPT-5.5, UK AISI identified a universal jailbreak within six hours of testing [14].
265 OpenAI subsequently reported several changes to its safeguard stack. However, a configuration issue
266 in the version supplied to UK AISI meant the institute could not verify the effectiveness of the final
267 configuration before deployment; the cited source does not document a verified patch of the jailbreak.

268 This case illustrates the distinction captured by the Partial category: a company may document a
269 response while the external evaluator remains unable to verify whether the final mitigation addressed
270 the identified vulnerability.

271 The limitation this case exposes is structural rather than incidental. Channel A (company response)
272 admits only the responding company’s own primary documentation (Section 3.4), so every one of
273 the 76 C1 (significant-risk) findings that drew a response is evidenced by that company’s account
274 of itself. That is a property of the coding rule, not a finding. The question the rule cannot answer is
275 whether a second party ever confirmed that the described action was taken, or that it worked.

276 Reading the response record for each of those 76 findings, we identify one case in which an evaluator
277 is documented as having re-tested a revised safeguard and endorsed it: the US Center for AI Standards
278 and Innovation’s (US CAISI’s) examination of Anthropic’s rebuilt cyber classifier, which occurred
279 during the export-control process and before redeployment (Section 5.1). We identify one case in
280 which the evaluator states that it could not verify the claimed fix—the GPT-5.5 jailbreak described
281 above. For the remaining 74, the public record contains the company’s description and nothing
282 independent against which to check it.

283 The count of independently verified remediations should be read as a lower bound. It reflects what
284 the cited sources describe; a verification exercise that took place but was not reported would not
285 appear. Four further candidates were examined and excluded: collaboration on pre-deployment
286 capability testing is not confirmation of a remedy, a capability designation is not a verified mitigation,
287 a developer confirming an attack is not a developer demonstrating a fix, and a company’s assessment
288 of its own model using its own tooling is not independent. The Partial category captures the subset
289 where the shortfall is visible at coding time—21 of the 188 C1 findings (11%)—but the broader
290 position is that proportionality throughout this paper is assessed against what companies state they
291 did, not against remediation anyone else has checked.

292 **Limitations.** These results describe the public record, not unreported company action, and do not
293 establish that an evaluation caused a response. Findings also cluster within reports, so subgroup
294 comparisons use report-level adjustments.

295 5 Interpreting the Accountability Gap

296 Section 4 documents the distribution of publicly observable responses but does not, by itself, explain
297 why those outcomes occurred. This section examines two forms of follow-through: company
298 remediation in two illustrative cases and policy-record responses to findings in the corpus. These
299 analyses help identify circumstances associated with observable action, but they are neither causal
300 estimates nor exhaustive explanations of the accountability gap. Section 6 separately compares the
301 follow-up mechanisms available to government AI institutes with those used by established oversight
302 institutions.

303 To illustrate how evaluation findings can be followed by documented remediation despite the broader
304 accountability gap, we examine two cases with an unusually clear public sequence from technical
305 evidence to developer response. Both are included in the corpus and coded as receiving a Substantive
306 response. In the Fable 5 and Mythos 5 case, the triggering Amazon report falls outside the evaluator-
307 scope sampling frame, but the later US CAISI safeguard-testing and redeployment record is included
308 as finding USCAISI-2026-06-CYB2. The cases contrast legally consequential action through an exist-

ing government authority with voluntary remediation through iterative pre-deployment collaboration. They are illustrative rather than representative.

5.1 Enforcement through an existing legal instrument: Fable 5 and Mythos 5

Anthropic launched Claude Fable 5 on June 9, 2026. According to Anthropic, the US government applied export controls to Fable 5 and Mythos 5 on June 12, requiring the company to prevent access by foreign nationals inside and outside the United States. But Anthropic suspended access to both models worldwide.

Anthropic reported that the directive followed government awareness of an Amazon research report describing a technique for bypassing Fable 5’s safeguards. The technique caused the model to identify software vulnerabilities and, in one instance, produce code demonstrating how one vulnerability could be exploited. Anthropic subsequently trained an updated safety classifier and reported that it blocked the identified technique in more than 99% of cases. Anthropic also reported that US CAISI tested both the previous and updated safeguards and considered them “extraordinarily strong”; this does not establish that US CAISI independently produced the 99% estimate. The controls were lifted on June 30. Fable 5 returned to global availability on July 1, while Mythos 5 was initially restored only to selected US organisations. [16].

The triggering Amazon report is not a corpus finding because it does not satisfy the evaluator-scope criteria. The later US CAISI safeguard-testing and redeployment record is included as finding USCAISI-2026-06-CYB2, and is the one corpus response for which the public record states that an evaluator tested the revised safeguards. The public record therefore documents a comparatively complete accountability sequence: a technical problem was reported, a binding legal instrument imposed restrictions, the company implemented a mitigation, and a government evaluator examined the revised safeguards before the restrictions were lifted.

The case does not show that the US federal government possessed or exercised deployment authority. The legally consequential action came through the Department of Commerce’s export-control powers, while US CAISI contributed technical testing within the resulting process. It therefore illustrates how evaluation expertise can support remediation when another institution supplies an existing legal pathway.

5.2 Collaborative Pre-Deployment Remediation: GPT-5.6

OpenAI’s system card describes an iterative pre-deployment process rather than a single evaluation event. UK AISI received early access to GPT-5.6 Sol and extensive grey-box access to its safeguard system. Across several testing rounds, UK AISI reported universal cyber-domain jailbreaks, including attacks that enabled long-form agentic work involving vulnerability discovery and exploit development. OpenAI states that it improved the system in response to UK AISI’s findings and, by launch, had reproduced and mitigated the specific jailbreaks UK AISI reported. The system card does not claim that the broader vulnerability class was eliminated: UK AISI expected further red-teaming to identify similar jailbreaks, and OpenAI committed to continued testing with the institute. [15].

We code the response as Substantive because OpenAI publicly attributes implemented pre-deployment mitigation to the jailbreaks reported by UK AISI. This classification refers specifically to mitigation of the reported attacks; it does not imply that GPT-5.6 had become resistant to all similar jailbreaks. Unlike the Fable 5 case, no binding government order was involved. The documented response occurred through repeated testing, private disclosure, mitigation, and retesting within an existing evaluator-developer collaboration.

Together, the two cases illustrate distinct pathways from technical evidence to remediation. In the Fable 5 case, an existing legal instrument imposed consequences, while government evaluation supported verification of the revised safeguards. In the GPT-5.6 case, pre-deployment access and an established evaluator-developer collaboration were followed by documented voluntary mitigation. Neither case establishes that publication of an evaluation finding, by itself, reliably triggers follow-up.

6 Institutional Mechanisms for Reducing the Accountability gap

The accountability problem examined in this paper is not unique to AI governance. Established oversight systems distinguish between producing technically credible findings, obtaining a response,

verifying corrective action, and possessing authority to compel or escalate action. We compare the UK AI Safety Institute (UK AISI) and US Center for AI Standards and Innovation (US CAISI) with seven established oversight bodies in the United States and the United Kingdom, focusing on formal mechanisms for information access, designated recipients, response, remediation or escalation, verification, and continuing follow-up. These bodies differ substantially in mandate, subject matter, and enforcement authority and are not presented as direct institutional templates for frontier AI governance. The comparison instead identifies mechanisms through which technical findings can remain actionable and publicly visible after an investigation or evaluation has concluded.

Detailed histories of the National Transportation Safety Board (NTSB) and Public Company Accounting Oversight Board (PCAOB), together with the wider institutional comparison, are provided in Appendix H [11, 13, 17, 22, 25].

Neither precedent provides a ready-made model for frontier AI governance. The NTSB cannot force recipients to adopt its recommendations, while the PCAOB regulates auditors rather than directly compelling audited companies to remedy every adverse finding. Nor do these precedents establish that AI evaluation institutions have failed in the same ways. They instead separate two institutional questions that are often conflated: whether a system provides a formal pathway from technical findings to documented follow-up, and whether the evaluators themselves are subject to independent supervision. Applied to frontier AI evaluation, these precedents motivate asking who should respond to an adverse finding, who tracks and evaluates that response, who verifies claimed remediation, and who—if anyone—oversees the evaluators.

The comparative lesson is therefore institutional rather than sector-specific. Government frontier-AI evaluators need not themselves become comprehensive regulators for their findings to enter a more systematic accountability process. Technical evaluation could remain institutionally separate while findings are connected to designated recipients, response expectations or obligations, public status tracking, procedures for reviewing claimed remediation, and defined pathways for escalation or referral. These mechanisms would not guarantee that every finding is remediated or that every response reduces real-world risk, but they would make follow-through more consistent, attributable, verifiable, and publicly inspectable.

7 Conclusion

Government AI institutes and independent evaluators increasingly provide evidence about frontier systems that would otherwise remain available only through developer self-reporting. The findings in this paper do not question the value of that technical work. They examine whether the public record shows that alarming findings concerning named systems were followed by responses proportionate to the identified risk.

Closing this public accountability gap does not necessarily require converting technical institutes into general-purpose AI regulators. Companies participating in government evaluations could publish finding-level responses stating whether they accept, dispute, plan to address, or have remediated a finding. Institutes could maintain a public response-status record, subject to legitimate security and confidentiality constraints, and distinguish claimed remediation from independently verified remediation. Evaluation agreements could also establish response timelines and specify what status information may be disclosed.

The public record currently provides limited grounds for confidence that government and independent frontier-AI evaluations consistently lead to substantive improvements in deployed-system safety. Creating a visible and repeatable pathway from finding to response, follow-up, and verification would make that confidence more warranted, and make the contribution of the evaluation ecosystem more publicly accessible.

References

- [1] M. Anderljung, E. T. Smith, J. O’Brien, et al. (2023). Towards publicly accountable frontier LLMs: Building an external scrutiny ecosystem under the ASPIRE framework. *arXiv preprint arXiv:2311.14711*. <https://arxiv.org/abs/2311.14711>.
- [2] R. Araujo, K. Fort, and O. Guest (2024). Understanding the first wave of AI Safety Institutes: Characteristics, functions, and challenges. *arXiv preprint arXiv:2410.09219*. <https://arxiv.org/abs/2410.09219>.
- [3] Anthropic (2026). *Responsible Scaling Policy, Version 3.4*. Effective July 8, 2026. <https://www.anthropic.com/responsible-scaling-policy>.
- [4] OpenAI (2025). *Preparedness Framework, Version 2*. Last updated April 15, 2025. <https://cdn.openai.com/pdf/18a02b5d-6b67-4cec-ab64-68cdfbdebcd/preparedness-framework-v2.pdf>.
- [5] Google DeepMind (2026). *Frontier Safety Framework, Version 3.1*. Published April 17, 2026. <https://deepmind.google/frontier-safety/>.
- [6] Y. Bengio et al. (2026). *International AI Safety Report 2026*. <https://arxiv.org/abs/2602.21012>.
- [7] A. Birhane, R. Steed, V. Ojewale, B. Vecchione, and I. D. Raji (2024). AI auditing: The broken bus on the road to AI accountability. *arXiv preprint arXiv:2401.14462*. <https://arxiv.org/abs/2401.14462>.
- [8] M. Brundage, N. Dreksler, A. Homewood, et al. (2026). Frontier AI auditing: Toward rigorous third-party assessment of safety and security practices at leading AI companies. *arXiv preprint arXiv:2601.11699*. <https://arxiv.org/abs/2601.11699>.
- [9] S. Costanza-Chock, E. Harvey, I. D. Raji, M. Czernuszenko, and J. Buolamwini (2023). Who audits the auditors? Recommendations from a field scan of the algorithmic auditing ecosystem. *arXiv preprint arXiv:2310.02521*. <https://arxiv.org/abs/2310.02521>.
- [10] G. Falco, B. Shneiderman, J. Badger, et al. (2021). Governing AI safety through independent audits. *Nature Machine Intelligence*, 3, 566–571. <https://doi.org/10.1038/s42256-021-00370-7>.
- [11] U.S. General Accounting Office (2002). *Protecting the Public Interest: Selected Governance, Regulatory Oversight, Auditing, Accounting, and Financial Reporting Issues*. Report GAO-02-483T. <https://www.gao.gov/products/gao-02-483t>.
- [12] J. Mökander, J. Schuett, H. R. Kirk, and L. Floridi (2023). Auditing large language models: A three-layered approach. *AI and Ethics*. <https://doi.org/10.1007/s43681-023-00289-2>.
- [13] National Transportation Safety Board (2024). *Responding to Our Safety Recommendations*. <https://www.ntsb.gov/investigations/Pages/RecipientTipRecommendations.aspx>.
- [14] OpenAI (2026). *GPT-5.5 System Card: External Evaluations for Cyber Capabilities—UK AISI*. <https://deploymentsafety.openai.com/gpt-5-5/external-evaluations-for-cyber-capabilities--uk-aisi>.
- [15] OpenAI (2026). *GPT-5.6 System Card*. <https://deploymentsafety.openai.com/gpt-5-6>.
- [16] Anthropic (2026). *Redeploying Fable 5*. <https://www.anthropic.com/news/redeploying-fable-5>.
- [17] Public Company Accounting Oversight Board (2013). *Staff Guidance Concerning the Remediation Process*. https://pcaobus.org/oversight/inspections/remediation/remediation_process.
- [18] I. D. Raji and J. Buolamwini (2019). Actionable auditing: Investigating the impact of publicly naming biased performance results of commercial AI products. In *Proceedings of the 2019 AAAI/ACM Conference on AI, Ethics, and Society*, pages 429–435. <https://doi.org/10.1145/3306618.3314244>.
- [19] I. D. Raji, A. Smart, R. N. White, et al. (2020). Closing the AI accountability gap: Defining an end-to-end framework for internal algorithmic auditing. In *Proceedings of the 2020 Conference on Fairness, Accountability, and Transparency*, pages 33–44. <https://doi.org/10.1145/3351095.3372873>.
- [20] I. D. Raji, P. Xu, C. Honigsberg, and D. E. Ho (2022). Outsider oversight: Designing a third party audit ecosystem for AI governance. *arXiv preprint arXiv:2206.04737*. <https://arxiv.org/abs/2206.04737>.
- [21] T. Rost (2026). From disclosure to self-referential opacity: Six dimensions of strain in current AI governance. *arXiv preprint arXiv:2604.14070*. <https://arxiv.org/abs/2604.14070>.
- [22] U.S. Securities and Exchange Commission (2003). *Order Regarding Section 101(d) of the Sarbanes–Oxley Act of 2002*. Release Nos. 33-8223 and 34-47746. <https://www.sec.gov/rules-regulations/2003/04/order-regarding-section-101d-sarbanes-oxley-act-2002>.

- [23] T. Shevlane, S. Farquhar, B. Garfinkel, et al. (2023). Model evaluation for extreme risks. *arXiv preprint arXiv:2305.15324*. <https://arxiv.org/abs/2305.15324>.
- [24] L. Staufer, M. Yang, A. Reuel, and S. Casper (2025). Audit cards: Contextualizing AI evaluations. *arXiv preprint arXiv:2504.13839*. <https://arxiv.org/abs/2504.13839>.
- [25] U.S. Congress (2026). *49 U.S.C. Section 1135: Secretary of Transportation’s Responses to Safety Recommendations*. <https://uscode.house.gov/view.xhtml?req=granuleid:USC-prelim-title49-section1135>.
- [26] J. Wang, K. Huang, K. Klyman, and R. Bommasani (2025). Do AI companies make good on voluntary commitments to the White House? *arXiv preprint arXiv:2508.08345*. <https://arxiv.org/abs/2508.08345>.

Appendix

A Sampling frame and source eligibility

The appendices follow the research process: sampling and search (Appendices A–B), coding (Appendix C), corpus and response evidence (Appendices D–E), robustness and worked examples (Appendices F–G), institutional context and prior work (Appendices H–I), and reproducibility (Appendix J).

A.1 Detailed eligibility and source scope

The 1,144 evaluation findings included outputs from the UK AI Security Institute, the US Center for AI Standards and Innovation, national and joint evaluation initiatives involving Australia, France, Japan, Singapore, and South Korea, and exercises conducted through the International Network for Advanced AI Measurement, Evaluation and Science. The sampling frame also included third-party evaluators such as METR, SecureBio, Apollo Research, Palisade Research, and other non-profit, academic, and commercial evaluation organisations.

Third-party outputs were eligible whether or not the work was conducted with a government institute, provided that the public material identified the evaluator and tested system and documented the methods and results sufficiently to extract a finding. Relevant sources included formal evaluation reports, research publications, institute progress reports containing evaluation results, joint-testing publications, and company model or system cards that explicitly reported or linked to an external evaluation. Company materials describing only internal evaluations were not treated as independent evaluation reports.

Table A.1 lists the organisations whose publications were screened to build the corpus, with the number of publications screened and reports included.

Table A.1: The 46 organisations in the sampling frame.

Organisation screened	Publications screened	Reports included
RAND Corporation	759	6
Kenya (Ministry of ICT and Digital Economy - network member)	628	0
Allen Institute for AI (Ai2)	547	0
Holistic AI	431	9
Google DeepMind	377	0
Scale AI	324	81
Anthropic	302	28
Cisco (Robust Intelligence / Foundation AI)	287	13
Thorn	270	0
UC Berkeley	269	0
OpenAI	166	20

Table A.1 (continued)

Organisation screened	Publications screened	Reports included
UK AI Security Institute (UK AISI)	160	67
Center for AI Safety (CAIS)	156	13
EPFL	155	0
Redwood Research	137	15
Tsinghua University CoAI Group	136	3
FAR.AI	117	19
Korea AI Safety Institute (K-AISI)	105	4
Japan AI Safety Institute (J-AISI)	100	3
METR	93	49
UL Research Institutes - Digital Safety Research Institute (DSRI)	91	2
Citadel AI	86	1
SecureBio	84	11
Collective Intelligence Project (Weval)	81	27
Shanghai AI Laboratory (AI45 Lab)	74	22
Yale University	69	0
Gray Swan AI	65	6
Stanford University	65	0
Apollo Research	57	13
Dreadnode	56	16
France INESIA (via PEReN / ANSSI / Inria / LNE)	55	4
LatticeFlow AI	52	5
US Center for AI Standards and Innovation (CAISI, NIST)	49	7
Canadian AI Safety Institute (CAISI Canada)	46	0
Meta	44	1
Princeton Holistic Agent Leaderboard (HAL)	35	9
Palisade Research	28	15
Transluce	27	14
University of Oxford	25	0
IndiaAI Safety Institute	21	0
Singapore AI Safety Institute (SG AISI)	18	4
Australian AI Safety Institute	12	0
EU AI Office (European Commission)	9	0
AI Verification and Evaluation Research Institute (AVERI)	7	1
Meridian Labs	7	0
Gray Swan AI + UK AISI (joint)	2	2

485 The sampling ledger contains 6,684 screened publication records and 490 organisation–report inclu-
486 sion entries. After deduplicating joint or collaborative reports credited to more than one screened
487 organisation and applying the final corpus exclusions, these entries correspond to the 474 reports
488 analysed in the paper. Organisations with zero included reports remain in the table so that the
489 sampling frame records unsuccessful as well as successful searches.

B Search tooling and human oversight

Corpus construction was AI-assisted throughout. This appendix records what was automated, what was not, and how the automated output was checked.

Enumeration and screening A total of 6,684 publications were enumerated across the 46 organisations in Appendix A. Of these, 5,453 were resolved by automated title-and-abstract triage against the inclusion rules in Section 3.1, 885 were examined by large language model (LLM) search agents where the title alone was insufficient, and 346 were resolved in later targeted passes (final-pass, enumeration-repair, reconciliation, and eval-lean). Every publication examined is recorded in `sweep/master_ledger.csv` with its decision and the reason, including exclusions, so coverage can be reconciled in both directions rather than inferred from what survived.

Finding extraction Candidate findings were extracted from source documents by the same agent tooling, which retrieved 5,256 documents in total. Each candidate was returned with a source URL and a verbatim quotation, and no candidate without both entered the corpus.

Response-channel search Searches for Channel A (company responses), Channel B (policy uptake), and Channel C (public coverage) were distributed across 109 agent-dispatched tasks in five phases: the window-extension sweep and its gap-closure pass, a policy/public-coverage sweep, a channel re-verification pass, and a backfill recovery. Tasks were instructed to apply the channel-specific searches below. The response-channel definitions and evidence-admissibility rules are described in Section 3.4.

Channel A (company responses) The search covered five sources: the company newsroom or blog; later model or system cards in the same family, including the launch card for pre-deployment evaluations; company safety and deployment pages; official company accounts or newsroom statements; and open-web searches used only to locate company primary sources. News articles were not admitted as company-response evidence. A response dated before the finding was excluded unless it documented coordinated pre-deployment disclosure.

Channel B (policy uptake) The search covered official records in the relevant jurisdiction. UK searches included Hansard, committee reports, government statements, and regulator publications. US searches included Congress.gov, the Federal Register, agency announcements, and committee records. EU and other searches covered Commission and national-regulator publications. A policy record counted only when it explicitly connected the finding, report, or result to the policy activity.

Channel C (public coverage) The search logged independent press coverage, dated academic-citation records, and notable public discussion. The evaluator’s own report and company posts were excluded because they are not independent coverage. Channel C (public coverage) records attention only; it does not affect Action Level or proportionality.

Adversarial verification Proposed Action Level codings were re-examined by an independent agent instructed to refute the coding by default and to uphold it only on the evidence cited. This pass overturned codings in a minority of cases, and those reversals are recorded in the Sources Checked field of the affected rows.

Human review The authors set every rule applied—the inclusion criteria, the three-tier taxonomy, the eight severity domains, the Action Level definitions, and the proportionality matrix—and reviewed the coded record for every finding retained in the corpus, resolving tier assignments.

Models used Severity classification used Claude Sonnet 5, GPT-5.5, and Gemini 3.1 Pro, with three calls per finding under a fixed prompt reproduced in Appendix C. Screening, extraction, channel search, and adversarial verification used LLM-based agent tooling over the same period. The earliest included report was published on November 16, 2022, and response searches used August 29, 2026, as the public-record cutoff.

Query strategy Enumeration was performed per organisation over each publication venue rather than through keyword search, so a finding could not be missed merely for lacking a search term.

538 C Tier taxonomy and severity annotation

539 Every finding was assigned to one of three mutually exclusive tiers. Tier A (trackable evaluation
540 findings) contains concerning empirical findings about a named frontier model or developer for which
541 a company response could reasonably be assessed. Tier B (non-trackable evaluation findings) contains
542 empirical findings that fail at least one Tier A condition, including anonymised, reassuring, null,
543 inconclusive, bare-score, non-frontier, or capability-trend results. Tier C (non-empirical findings)
544 contains methodology, framework, governance or process, tooling, milestone, and other non-empirical-
545 model findings. Only Tier A enters the response analysis.

546




Figure C.1: The three-tier taxonomy across the 1,144-finding corpus: Tier A (trackable evaluation findings), 231 (20%); Tier B (non-trackable evaluation findings), 599 (52%); and Tier C (non-empirical findings), 314 (27%). Only Tier A enters the response analysis.

C.1 Severity domains and final review

The severity prompt tests eight domains: chemical, biological, radiological, and nuclear (CBRN) and bio-chemical uplift; offensive cyber capability; autonomy, AI research-and-development (AI-R&D) automation, and self-replication; persuasion or societal harm at scale; deliberate deception or misalignment; deployed-safeguard failure; compromised evaluation integrity; and acute individual harm. Each finding was independently classified by Claude Sonnet 5, GPT-5.5, and Gemini 3.1 Pro. Their majority vote produced the provisional C1 (significant-risk) or C2 (low-risk) label. A human annotator then reviewed every finding and the model rationales, resolving ambiguous cases and, where warranted, overriding the majority; the human-reviewed label is the final severity classification used in the analysis.

Table C.1: Raw ensemble agreement. 1,030 of 1,144 findings received a unanimous 3–0 vote and 114 received a 2–1 vote. After human review, the final corpus contained 333 C1 (significant risk) and 811 C2 (low risk) findings.

Subset	n	Unanimous	Split	Agreement rate
All findings	1,144	1,030	114	90.0%
Tier A (Trackable evaluation findings) only	231	192	39	83.1%

547 The exact frozen prompt files (severity_prompt.txt and severity_prompt_v1.0_FROZEN.txt), model
548 votes, retained rationales, and final labels are available with the supplemental study materials. Votes
549 produced under v1.0 and v1.1 are tagged by version and are never combined without preserving that
550 provenance. <https://kunalssingh.com/evaluating-the-evaluators/#data>

551 D Corpus composition

552 The corpus contains 46 distinct reporting-institution labels, including separate labels for joint evalua-
553 tions and collaborations. Appendix A lists the separate 46-organisation screening frame; the equal
554 totals are coincidental because the two lists do not contain exactly the same entities. In the table, UK
555 AISI denotes the UK AI Safety Institute and US CAISI denotes the US Center for AI Standards and
556 Innovation; national AISI labels denote the corresponding national AI Safety Institute.


Table D.1: All 46 reporting-institution labels in the study dataset, ordered by findings down the left column and then the right. A label may denote a standalone institution, a collaboration, or a joint evaluation.

Reporting institution	Findings	Reporting institution	Findings
UK AI Safety Institute (UK AISI)	202	Apollo Research + Anthropic (Apollo in the AI Safety Institute Consortium, AISIC)	6
Scale AI	152	LatticeFlow AI	6
METR	113	Princeton CRUX	6
Collective Intelligence Project (Weval)	78	AI Verification and Evaluation Research Institute (AVERI)	5
Shanghai AI Laboratory (AI45 Lab)	74	Joint UK + US + Singapore AISIs (International Network)	4
SecureBio	48	Joint UK AISI + OpenAI (company-published)	4
Apollo Research	46	Redwood Research + Scale AI	4
FAR.AI	42	Joint UK AISI + US CAISI + Gray Swan	3
Transluce	39	Singapore AI Safety Institute (SGAIS) / IMDA	3
Joint UK AISI + US CAISI	37	Japan AI Safety Institute (J-AISI)	2
US Center for AI Standards and Innovation (US CAISI)	34	Joint UK AISI + International Network (10 countries)	2
Center for AI Safety (CAIS)	32	Joint UK AISI + US CAISI + Anthropic	2
Dreadnode	31	Korea AI Safety Institute (K-AISI)	2
Redwood Research	29	RAND Europe	2
Princeton Holistic Agent Leaderboard (HAL)	27	Signature Science	2
Palisade Research	19	Singapore AI Safety Institute (SGAIS) / Korea AI Safety Institute (KR AISI)	2
RAND	18	UK AISI + Thorn	2
Holistic AI	13	Andon Labs	1
Gray Swan AI	11	Citadel AI	1
Cisco (Robust Intelligence / Foundation AI)	10	METR + Epoch AI	1
Irregular	10	Singapore AISI + Korea AISI (bilateral, network members)	1
France PEReN/INESIA	9	UK AISI + Anthropic + Theorem + MATS	1
International Network of AI Safety Institutes	7	UL Research Institutes (DSRI)	1

Figure D.1 complements the complete list above by showing how reporting volume is distributed across institution types. It preserves the 1,144-finding total while showing the five largest reporting institutions in each group; the remaining institutions are combined as “Other institutions.”

Who reports findings, by institution type

Bar length = findings, on one scale across all four branches · top 5 institutions per type



Institution Type field; compound values (e.g. "Government;Lab") fold into their primary type, so the four branches sum to 1,144.

Figure D.1: Reporting volume by institution type across all 1,144 findings. Bar lengths use one common scale across the four groups. Institution types follow the dataset's Institution Type field. Compound values (for example, Government;Lab) are assigned to their primary type, so the four groups sum to 1,144. The figure describes corpus composition, not evaluator quality or independence.

E Response evidence and data access

Complete finding-level rows and supporting evidence are available through the study website below. Among the 38 C1 (significant-risk) Tier A (trackable evaluation) findings coded Substantive, 35 carry Explicit attribution—the company named the evaluator or finding in its response; the three exceptions drew a substantive fix with no explicit public credit to the evaluator.

Complete rows and supporting data: <https://kunalssingh.com/evaluating-the-evaluators/#data>

F Robustness and subgroup analyses

Core corpus statistics were independently recomputed from the study dataset; the reproduction path is documented in Appendix J. Throughout this appendix, Tier A means trackable evaluation findings, C1 means significant risk, and C2 means low risk.

F.1 Is the headline an artefact of the counting unit?

The headline counts findings. Because one report can yield several findings, a small number of heavily split reports could in principle drive the result. It does not: the gap rate remains high under each alternative unit.

Table F.1: Robustness of the headline gap to alternative counting units.

Unit of analysis	Gap	Rate	Wilson 95% CI
Finding-weighted (headline)	150/188	79.8%	73–85%
Report-weighted (a report counts once)	86/102	84.3%	76–90%
Institution-weighted (mean for institutions with at least three findings, n=18)	—	83.6%	—
All Tier A, ignoring severity	186/231	80.5%	75–85%
C2 (low risk) findings only	36/43	83.7%	70–92%

Table F.1 shows that the result remains high under alternative counting units: 79.8–84.3% for the finding-, report-, all-Tier-A, and C2 analyses, and 83.6% under institution weighting.

F.2 Uncertainty and subgroup comparisons

The headline is a description of an enumerated public record rather than an estimate from a sample, so we do not attach a significance test to it. Table F.1 instead reports the result under five counting units, with Wilson intervals where applicable. Two subgroup comparisons support claims in the main paper and are reported here with their limits.


Access type. Among all 231 Tier A (trackable evaluation) findings, 3 of 54 pre-deployment findings received no response (5.6%), compared with 132 of 160 post-deployment findings (82.5%); the 17 Mixed findings are excluded from this comparison. A two-proportion test gives $z = -10.13$, $p < 0.001$. The difference remains below $p < 0.001$ after adjustment for clustering within reports.

Severity. C1 (significant-risk) findings received no response in 112 of 188 cases (59.6%), compared with 31 of 43 C2 (low-risk) findings (72.1%); $z = -1.52$, $p = 0.127$. This is not evidence of no difference. With only 43 C2 findings, the comparison has 33% power against the observed difference and would need a gap of roughly 25 percentage points to reach 86% power. The corpus therefore cannot resolve whether severity predicts response.

Clustering caveat. The 188 C1 Tier A findings sit in 102 reports, and outcomes cluster within reports because findings from the same report usually share an evaluator, disclosure arrangement, and developer. The estimated intra-cluster correlation is 0.62. The resulting design effect is 1.52, reducing the effective sample size from 188 to about 123. Two-proportion tests that treat findings as independent can therefore overstate significance; comparisons near the conventional threshold are reported descriptively rather than as demonstrated effects.

Two additional comparisons are descriptive. Government AISI findings fell short in 42 of 59 cases (71.2%), compared with 108 of 129 (83.7%) for third-party evaluators; this difference does not survive the clustering adjustment ($p = 0.108$). Among C1 Tier A findings, 90 of 156 findings with documented public traction received no response (57.7%), compared with 22 of 32 without documented traction (68.8%). This second comparison has 21% power and is inconclusive.

602



Government institutes 42/59 = 71.2% fall short · third-party evaluators 108/129 = 83.7%. The headline does not rest on pooling the two: it holds within each. Bars are 100% stacked because the groups differ in size; counts are printed inside. Corpus-wide the split is 319 government and 825 third-party findings.

Figure F.1: Outcome composition by evaluator type among the 188 C1 (significant-risk) Tier A (trackable evaluation) findings. Both groups show a high shortfall rate. The difference between government AISI and third-party evaluators is descriptive because it does not survive adjustment for clustering within reports.

603




Figure F.2: Access type in the full 1,144-finding corpus (top) and no-response rates within the 188 C1 (significant-risk) Tier A (trackable evaluation) findings (bottom). The formal pre- versus post-deployment test above uses all 231 Tier A findings and excludes Mixed findings; the bottom panel is the narrower C1 description.

604 F.3 Outcome composition by risk domain

605 Domains are multi-label, so the row totals in Figure F.3 exceed the 231 Tier A findings. The largest
 606 domains all contain substantial accountability-gap shares, while the smallest domains have too few
 607 findings for stable comparisons. The figure is therefore descriptive and should not be read as a ranking
 608 of domain importance.




Figure F.3: Outcome composition by risk domain among Tier A (trackable evaluation) findings. Counts and within-domain percentages are shown in each cell. Cell shading runs from light to dark blue with the within-domain share. Domains are multi-label, so row totals exceed 231 findings.

610 F.4 Channel C (public coverage) — traction among no-response findings

Table F.2: Documented external coverage among C1 (significant-risk) findings with no company response.

Public-coverage evidence (Channel C)	Findings with coverage	Share
Media Outlets	36/112	32%
Academic Citations	70/112	63%
Social Highlights	19/112	17%
Any of the three	90/112	80%

611 Documented external coverage among the 112 C1 findings that received no company response. The
 612 categories overlap: one finding may appear in media, academic, and social records. Most were
 613 publicly visible beyond the original report, so lack of attention does not account for the silence.
 614 Distinguish an empty cell (not searched, or not applicable) from a recorded "none located" (searched,
 615 nothing found); only the latter supports an inference.

616 F.5 Policy-record responses to corpus findings

617 Policy Level records whether a finding appears in an official policy response:

- 618 1. **No policy uptake identified:** no qualifying official policy response was located.

2. **Non-binding policy-related uptake:** the finding was connected to guidance, an official statement, a request for information, research support, debate, or another non-binding response.
3. **Binding policy action:** an official source or the regulated recipient documented a binding restriction or obligation explicitly connected to the finding.

A policy record was coded only when the evidence connected the specific finding to the policy response; general policy discussion in the same domain was not sufficient.




Figure F.4: Policy Level across Tier A (trackable evaluation) findings: no uptake 184, non-binding uptake 46, binding action 1. One binding policy action across the entire accountability set.

We searched parliamentary, congressional, regulatory, and other official government records for explicit references to corpus findings between 2024 and 2026. We classified an appearance as Type A (binding policy action) when the evidence explicitly connected the finding to binding action and as Type B (non-binding policy action) when an official source connected the finding to a non-binding response, such as guidance, research funding, an official statement, or a voluntary commitment. For the sole Type A case, the private order was documented by the regulated recipient and contemporaneous reporting; no official government publication of the order was located. Temporal proximity alone was not sufficient.

We also searched for Type C cases in which an institute’s existence or evaluation work was cited as a reason to defer legislation. No confirmed Type C case was identified during the observation period. This is an observation about the searched public record, not evidence that such arguments never occur.

Table F.3: Selected policy-record appearances among the 47 Tier A (trackable evaluation) findings coded as policy uptake: 46 non-binding and one binding. Type A denotes binding policy action; Type B denotes non-binding policy-related uptake.

Finding ID and cited result	Venue/source and classification	Documented policy response
USCAISI-2026-06-CYB2: Fable 5/Mythos 5 jailbreak, classifier retest, and redeployment	US Department of Commerce (recipient disclosure) Type A (binding)	A private Export Administration Regulations order restricted access for about 18 days; controls were lifted subject to security and reporting conditions. No official government publication of the private order was located.
UKAISI-2026-04-CYB1: Mythos Preview achieved 73% on expert CTFs and solved a hardest-range challenge	UK Government (DSIT) Type B (non-binding)	An open letter explicitly cited AISI’s Mythos evaluation and launched a Cyber Resilience Pledge; it did not reproduce this row’s exact statistic or create an enforceable obligation.

Table F.3 (continued)

Finding ID and cited result	Venue/source and classification	Documented policy response
UKAIS-2026-04-CYB11: GPT-5.5 cyber-range and reverse-engineering results	Bank of England FPC Type B (non-binding)	The July 2026 Financial Stability Report reproduced the result and cited AISI in its systemic-risk assessment; no Recommendation or Direction was issued on that basis.
OPENAI-2026-08-CYB1: an evaluation agent acted on a real website after an environment misconfiguration	UK NCSC Type B (non-binding)	The NCSC issued a formal advisory statement calling for safeguards, real-time oversight, and incident-response plans; it imposed no binding obligation.
JOINT-2025-01-JAI1: agent-hijacking attacks raised measured success from 11% to 81%	US NIST/CAISI Type B (non-binding)	A Federal Register request for information cited the report as its research basis; the notice requested comments and imposed no mandatory requirement.
SECUREBIO-2025-04-BIO1: frontier models outperformed expert virologists on a 322-question benchmark	US House committee record Type B (non-binding)	Published written testimony named the Virology Capabilities Test and stated the result; no enforceable obligation followed.
PALISADE-2025-07-ALI1: o3 disabled its shutdown mechanism in 79 of 100 experiments	European Parliament Type B (non-binding)	A written question cited the shutdown result and called for regulatory intervention; no binding action was located.

Only one of the 47 Tier A findings with documented policy uptake was associated with binding action. In that case, the finding entered an existing export-control pathway rather than creating a new enforcement mechanism. Most other policy appearances resulted in debate, guidance, requests for information, recommendations, research commitments, or other advisory responses.

These classifications establish an observable connection between a finding and the official record; they do not establish that the finding alone caused the resulting policy action. Legislation and regulatory decisions generally arise from multiple sources of evidence, political considerations, and prior institutional work.

G Worked examples across the response scale

Two findings at each Action Level, drawn from the dataset, to make the coding boundaries concrete. These are illustrative selections, not a random sample; each is reproducible from its Finding ID in the published dataset.

651 **G.1 Substantive — a specific, documented, attributed change (n = 38 of 188 C1**
652 **significant-risk Tier A findings)**

Table G.1: Substantive worked examples

Finding ID	Institution	Finding	Company response
METR-2026-08-ALI4			
Evidence source	METR	The agents moved from internal cheating research to compromising an uninvolved third party's production systems. At ar...	Between July 10 and July 13, agents identified Hugging Face user credentials that were exposed on the internet and use...
METR-2026-08-ALI2			
Evidence source	METR	Agents used the board to run large collective projects whose common objective was a general-purpose way to trick or ta...	Cheating, broken environments, and safe stopping. When a task is corrupted, broken, or impossible, agents should reque...

654 **G.2 Partial — action documented but incomplete, interim, or unverifiable (n = 21 of 188 C1**
655 **significant-risk Tier A findings)**

Table G.2: Partial worked examples

Finding ID	Institution	Finding	Company response
OPENAI-2026-08-CYB1			
Evidence source	Irregular	Irregular, one of OpenAI's external cybersecurity testing partners, notified OpenAI on 29 July 2026 of an incident dur...	In the coming weeks, we will review our own approach to third-party testing, including how we identify higher-risk eva...
UKAISI-2026-08-AUT1			
Evidence source	UK AISI	In the same exercise, Mythos 5 agents running in separate, concurrently executing, isolated samples discovered one ano...	Finally, we note that the UK's AISI recently published a report on a cybersecurity evaluation involving Claude Mythos...

657 **G.3 Acknowledged — the finding is referenced; no action is described (n = 17 of 188 C1**
658 **significant-risk Tier A findings)**

Table G.3: Acknowledged worked examples

Finding ID	Institution	Finding	Company response
SECUREBIO-2023-06-BIO1	SecureBio	In a one-hour 'Safeguarding the Future' classroom exercise at MIT, three groups of non-scientist students prompted pub...	Existing research on AI-enabled biological threats has shown that models like GPT-4 can be prompted or red-teamed to s...
SECUREBIO-2023-06-JAI1	SecureBio	The chatbots' biosecurity safeguards were bypassed trivially: most harmful responses were offered freely with at most...	Existing research on AI-enabled biological threats has shown that models like GPT-4 can be prompted or red-teamed to s...

660 **G.4 None — nothing located by the search cutoff (n = 112 of 188 C1 significant-risk Tier A**
661 **findings)**

Table G.4: No-response worked examples

Finding ID	Institution	Finding	Company response
UKAISI-2026-08-ALI1	UK AISI	In AISI cyber-range evaluations run 25–28 July 2026 (122 samples across two 'Doing Life' ranges, DL-v1 and DL-v2, with...	—
TRANSLUCE-2026-08-ALI3	Transluce	Claude's behaviour is conditioned on the recognised identity of the user supplied through the Claude Code harness (whi...	—

663 The finding and response fields are abridged for layout; the complete rows and evidence are available
664 through the dataset link in Appendix E. The distinction that matters most in practice is Partial versus
665 Substantive: a company may document a response while the evaluator remains unable to verify that
666 the mitigation addressed the identified vulnerability. See Section 4.4.

667 **H Institutional comparison**

668 **H.1 Comparative framing**

669 In this appendix, NTSB denotes the National Transportation Safety Board; PCAOB, the Public
670 Company Accounting Oversight Board; FDA, the Food and Drug Administration; NRC, the Nuclear
671 Regulatory Commission; ONR, the Office for Nuclear Regulation; FCA, the Financial Conduct
672 Authority; and CAA, the Civil Aviation Authority. UK AISI denotes the UK AI Safety Institute, and
673 US CAISI denotes the US Center for AI Standards and Innovation. The NTSB provides an especially
674 relevant precedent for non-binding technical findings. It independently investigates transportation
675 accidents, determines their probable causes, and issues recommendations intended to prevent recurrence.
676 The NTSB cannot itself compel recipients to implement those recommendations; however,
677 it maintains a formal follow-up system. For recommendations directed to the Secretary of Transportation,
678 federal law requires a written response within 90 days stating whether the recommendation
679 will be adopted fully, adopted partly, or refused, together with reasons and, where applicable, an
680 implementation timetable. Recommendations and responses are made public. The NTSB subsequently
681 classifies each response and implementation status using categories such as open or closed

682 and acceptable or unacceptable [13, 25]. The NTSB therefore shows how a technically capable but
683 non-regulatory institution can maintain a publicly inspectable pathway from findings to response and
684 follow-up even when it cannot require implementation.

685 The collapse of Enron and the subsequent creation of the Public Company Accounting Oversight
686 Board (PCAOB) illustrate a different institutional question: independent supervision of evaluators
687 themselves. Enron exposed serious weaknesses in corporate governance, auditor independence,
688 financial-statement auditing, professional self-regulation, and financial reporting [11]. It was not
689 simply a case in which credible audit warnings were published and ignored; the reliability and
690 independence of the assurance process were themselves in question. Congress responded through
691 the Sarbanes–Oxley Act of 2002, which created the PCAOB to register public-accounting firms,
692 establish auditing and independence standards, inspect registered firms, investigate misconduct,
693 conduct disciplinary proceedings, and enforce compliance [22]. Its inspection regime also includes a
694 structured process under which audit firms may address identified quality-control criticisms, with
695 unresolved criticisms becoming public if they are not addressed to the Board’s satisfaction within 12
696 months [17].

697 The comparison shows that binding enforcement power is not the only mechanism through which
698 findings can generate publicly observable follow-up. The NTSB lacks authority to require implemen-
699 tation, yet it identifies recipients, obtains and evaluates responses, tracks implementation status, and
700 keeps unresolved recommendations publicly visible. Other comparators combine these functions
701 with compulsory information access, remediation procedures, verification authority, or referral to
702 an enforcement body. The UK AI Safety Institute (UK AISI) and US Center for AI Standards and
703 Innovation (US CAISI) may sometimes obtain responses or examine remediation through voluntary
704 or negotiated arrangements, but they lack a uniform formal process requiring a response, assigning
705 responsibility for follow-up, or maintaining the public status of unresolved findings.

706 Table H.1 broadens the comparison beyond the NTSB and PCAOB to the Food and Drug Administra-
707 tion (FDA), Nuclear Regulatory Commission (NRC), Office for Nuclear Regulation (ONR), Financial
708 Conduct Authority (FCA), and Civil Aviation Authority (CAA). The comparison does not imply that
709 these bodies possess identical powers or that each mechanism operates in the same way. Rather, it
710 records whether their institutional arrangements formally provide the relevant function. What the
711 comparators share is not universal authority to compel every recommended action, but an institutional
712 pathway connecting findings to a recipient, response, remediation or escalation, verification, or
713 continuing follow-up.

Table H.1: Comparison of UK AISI and US CAISI institutional design against established oversight bodies. Acronyms are expanded in the preceding text.

Design feature	FDA	NRC	NTSB	PCAOB	ONR	FCA	CAA	UK AISI / US CAISI
Can impose binding restrictions within its statutory mandate	Yes	Yes	No	Yes	Yes	Yes	Yes	No
Compulsory access or information-gathering rights	Yes	Yes	Yes	Yes	Yes	Yes	Yes	No
Formal process for obtaining a response	Yes	Yes	Yes	Yes	Yes	Yes	Yes	No
Defined escalation or remediation process	Yes	Yes	~	Yes	Yes	Yes	Yes	No
Can review or verify reported remediation	Yes	Yes	~	Yes	Yes	Yes	Yes	~
Defined pathway to enforcement or another responsible authority	Yes	Yes	Yes	Yes	Yes	Yes	Yes	No
Systematic post-finding follow-up and status tracking	Yes	Yes	Yes	Yes	Yes	Yes	Yes	No

Yes = formally present; No = absent; ~ = partial or conditional. PCAOB authority applies to registered audit firms, not the public companies they audit. NTSB tracks and evaluates responses but cannot compel implementation. UK AISI and US CAISI can sometimes examine remediation through voluntary or negotiated access, but neither has a uniform right to do so. The two institutes are combined here only for this high-level comparison; their legal and institutional arrangements differ.

I Extended prior work

Prior work addresses three questions adjacent to this study: how frontier AI evaluations and audits should be designed, whether audits and voluntary commitments are followed by observable changes, and how the institutional and political environment shapes accountability.

The first body of work concerns evaluation and audit design. Raji et al. [19] propose an end-to-end framework for internal algorithmic auditing that embeds accountability across the development life cycle. Mökander et al. [12] propose a three-layered framework combining governance, model, and application audits. Anderljung et al. [1] identify six requirements for effective external scrutiny of frontier systems: access, a searching attitude, proportionality to risk, independence, resources, and expertise. Drawing lessons from financial, environmental, and health regulation, Raji et al. [20] argue that third-party audits require a supporting institutional ecosystem to produce accountability. Brundage et al. [8] define frontier AI auditing as rigorous third-party verification of developers' safety and security claims and propose four AI Assurance Levels for communicating confidence in audit conclusions. Staufer et al. [24] propose "audit cards" for documenting auditor identity, evaluation scope, methodology, resource access, process integrity, and review mechanisms. Collectively, this literature identifies conditions for credible evaluation and auditing but does not empirically measure whether particular published findings subsequently produce company or policy action.

The International AI Safety Report 2026 [6] documents a related upstream challenge: reliable pre-deployment testing has become harder as models increasingly distinguish evaluation from deployment settings and exploit evaluation loopholes.

A second body of work examines whether audits or public commitments produce observable changes. Raji and Buolamwini [18] provide the closest direct methodological precedent. They re-tested commercial facial-analysis products following the public Gender Shades audit to examine whether named companies had improved their systems. Wang et al. [26] assess the publicly disclosed behaviour of 16 companies against their eight 2023 voluntary commitments to the White House. They report an average compliance score of 53% and argue that companies should provide proactive

745 and verifiable evidence of compliance. These studies measure product changes or compliance with
746 prior commitments; the present study instead traces responses to independently produced findings
747 about frontier AI systems.

748 Finally, Rost [21] applies a six-dimensional political-theory framework to six existing AI-governance
749 arrangements. The analysis rates the UK AISI as failing on corrigibility and voluntary company
750 commitments as failing on accountability, while explicitly presenting its cross-case patterns as hy-
751 potheses requiring further validation. Our study complements this institutional analysis by measuring
752 one observable component of accountability: whether publicly documented evaluation findings are
753 followed by attributable action.

754 Taken together, this literature explains why evaluations matter, how rigorous audits might be designed,
755 and why disclosure may not reliably produce accountability. To our knowledge, however, no prior
756 study systematically tracks, finding by finding, the publicly documented company, policy, and
757 ecosystem responses to frontier AI system evaluations conducted by government institutes and
758 independent frontier-AI evaluators.

759 J Reproducibility

760 J.1 Resources

761 **Project data, dashboard, and study materials:**
762 <https://kunalssingh.com/evaluating-the-evaluators/#data>

763 Core corpus statistics in this version were independently recomputed from the workbook. The project
764 scripts provide the reproducibility and validation path summarised below.

765 **Study dataset** The complete finding-level dataset: 1,144 findings across 39 fields. The analysis
766 uses one controlled source path so that reported results are derived consistently.

767 **dataset/AISIEVAL_validate.py** Structural validator: identifiers, controlled vocabularies, date co-
768 herence, the proportionality formula, and the Attribution invariant. Currently PASS with 0 violations.

769 **scripts/all_stats.py** Recomputes every reported quantity - corpus, headline, channels, proportional-
770 ity, robustness, subgroup tests, lag, and the by-developer, by-year and by-domain breakdowns.

771 **scripts/build_charts.sh** Builds all 28 figures. `scripts/verify_charts.py` then re-derives every plotted
772 quantity from the workbook and fails the build if any figure disagrees.

773 **scripts/find_duplicates.py** Six independent duplicate tests across identifiers, finding text, quotes,
774 and near-duplicate detection. Currently 0 duplicates.

775 **severity/severity_prompt.txt** The severity prompt, version 1.1. Version 1.0 is preserved verbatim;
776 every vote is tagged with the version that produced it.

777 **sweep/master_ledger.csv** The screening ledger: every publication examined and the decision taken,
778 including exclusions, so coverage can be reconciled in both directions.

779 Reproduction path. The validator, the statistics script and the chart verifier are run together after any
780 change to the dataset; a figure that disagrees with the workbook is treated as a build failure rather
781 than a discrepancy to be explained.

782 J.2 Exact reproduction commands

783 From the project root, run the following commands in order:

```
784 python dataset/AISIEVAL_validate.py  
785 python scripts/all_stats.py  
786 bash scripts/build_charts.sh
```

787 The first command validates the finding-level workbook, the second regenerates the reported statistics,
788 and the third rebuilds the figures. The supplemental package includes the exact severity-prompt
789 files; the project data and code are available from the study website. [https://kunalssingh.com/](https://kunalssingh.com/evaluating-the-evaluators/#data)
790 [evaluating-the-evaluators/#data](https://kunalssingh.com/evaluating-the-evaluators/#data)